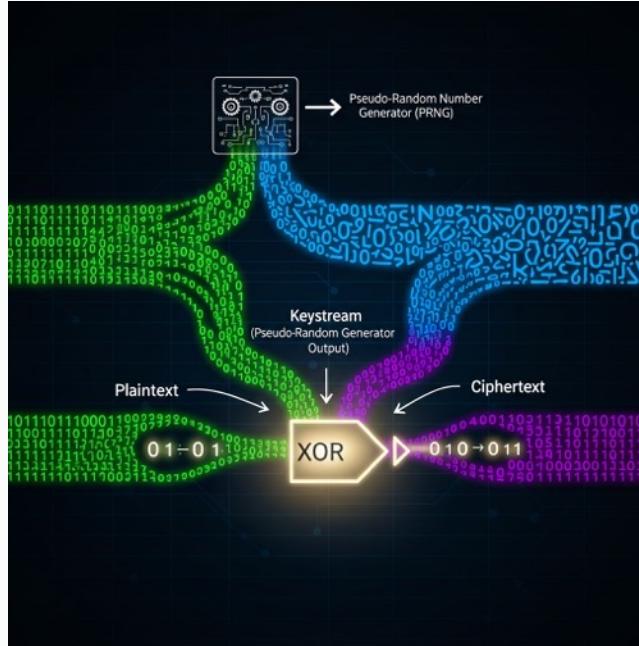


# Stream ciphers

old, but now back in  
the spotlight



# Secret key cryptography

Alice and Bob share

- A crypto protocol  $E$
- A secret (symmetric) key  $K$
- They communicate using  $E$  with key  $K$
- Adversary knows  $E$ , knows some exchanged messages but ignores  $K$

Two approaches:

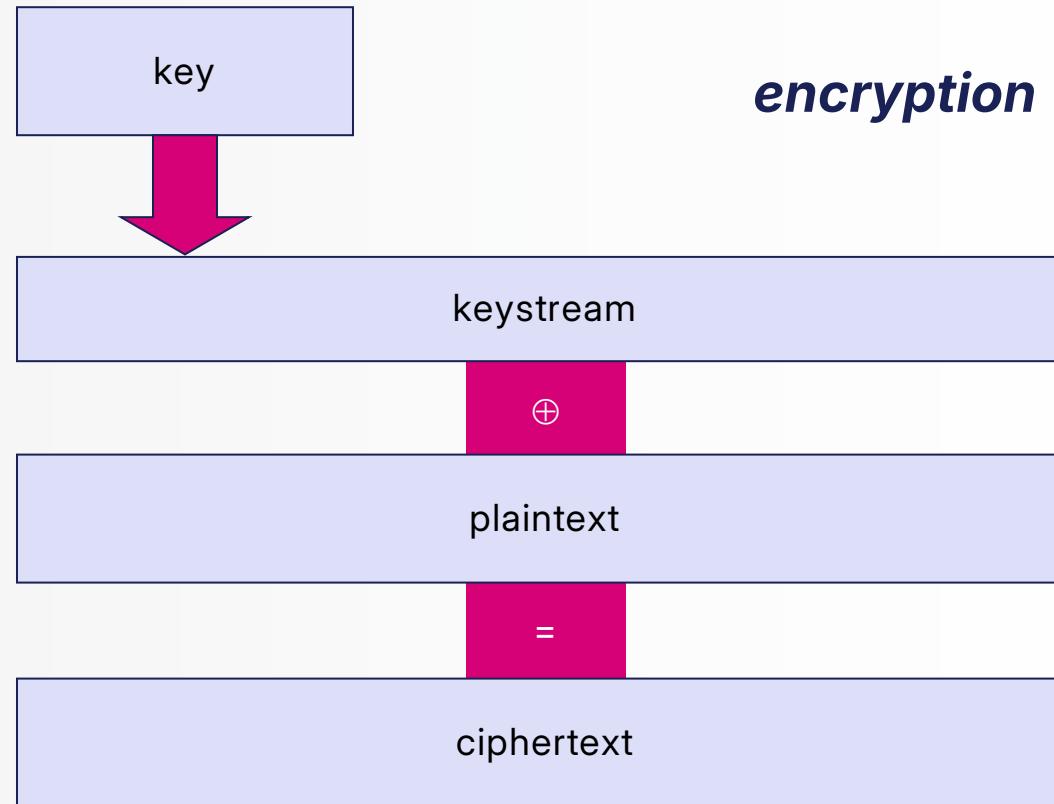
- Stream Ciphers
- Block ciphers

# Stream ciphers

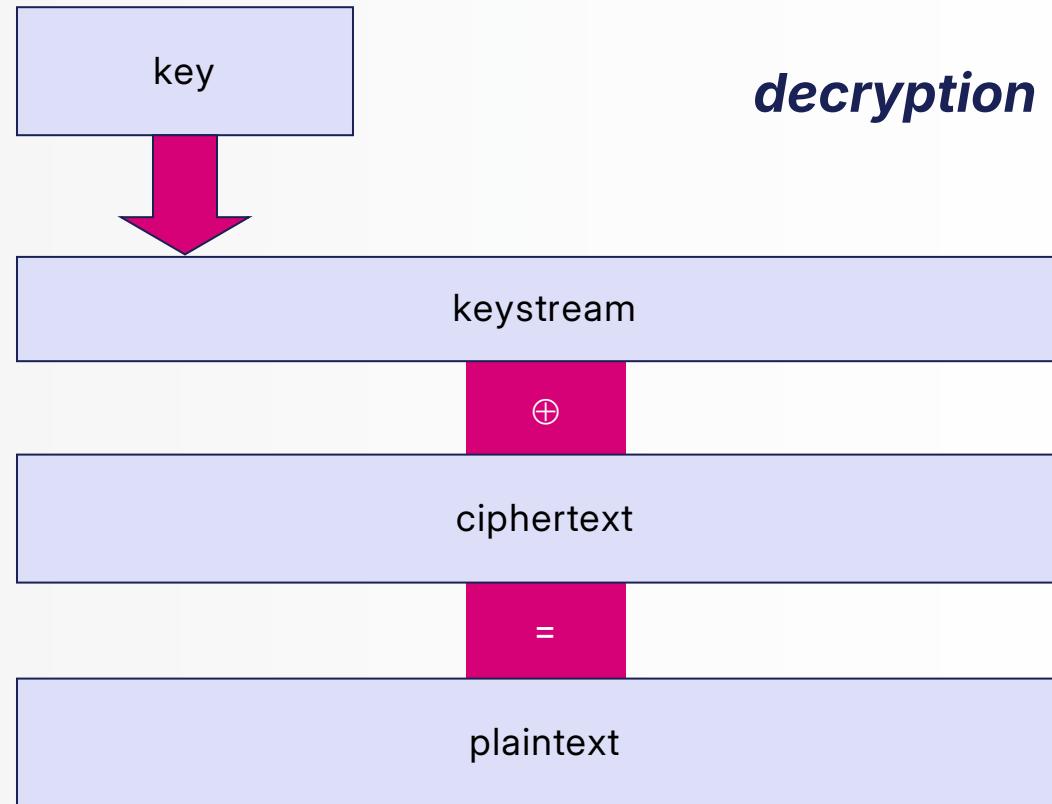
- define a secret key (*seed*)
- using the seed generate a byte stream (*keystream*):  
*i*-th byte is function of
  - only key (*synchronous* stream cipher), or
  - both key and first *i*-1 bytes of ciphertext  
*asynchronous* stream cipher)
- obtain ciphertext by bitwise XORing plaintext and keystream

$x$	$y$	$x \oplus y$
0	0	0
0	1	1
1	0	1
1	1	0

# Synchronous stream cipher



# Synchronous stream cipher



# Synchronous stream ciphers in practice

- Many ciphers before 1940
- Enigma - II world war (Germany)
- A5 – GSM (encryption cell phone-base station)
- WEP - used in Ethernet 802.11 (wireless)
- RC-4 (Ron's Code, used by WEP)

# A5/1 stream cipher

- Developed in 1987, used to ensure over-the-air communication privacy in the GSM cellular telephone standard
- A major dispute arose in the mid-1980s between NATO signal intelligence agencies about whether GSM encryption should be strong or intentionally weakened
- Used in both Europe and the United States. A related variant, A5/2, was a deliberately weakened version intended for export to certain regions
- The design was initially kept secret, but leaked in 1994 and fully reverse-engineered in 1999 by Marc Briceno
- Multiple serious vulnerabilities in the cipher were subsequently identified



# RC-4

- RC: Ron's Code
  - (Ron = Ronald Rivest, MIT, born in 1947 in NY state)
- Considered safe: 1987 - 1994 kept secret, after '94 extensively studied
  - insecure from 2004
- Good for exporting (complying with US restrictions)
- Easy to program, fast
- Very popular: Lotus Notes, SSL, Wep etc.
- RC4's weak key schedule can give rise to a variety of serious problems

# RC-4 weaknesses

- Biased output in early keystream bytes
- Vulnerable if keys are reused
- Broken in protocols like WEP and TLS (early versions)
- Today deprecated

# RC-4 properties

- variable key length (byte)
- synchronous
- starting from the key, it generates an apparently random permutation
- eventually the sequence will repeat
- however, long period  $> 10^{100}$
- very fast: 1 byte of output requires 8-16 instructions

# RC-4 initialization

Goal: generate a (pseudo)random permutation of the first 256 natural numbers

1.  $j=0$
2.  $S_0=0, S_1=1, \dots, S_{255}=255$
3. Assume a key of 256 bytes  $k_0, \dots, k_{255}$  (if the key is shorter, repeat)
4. for  $i=0$  to 255 do
  1.  $j = (j + S_i + k_i) \bmod 256$
  2. exchange  $S_i$  and  $S_j$

In this way we obtain a permutation of 0, 1, ..., 255, the resulting permutation is a function of the key

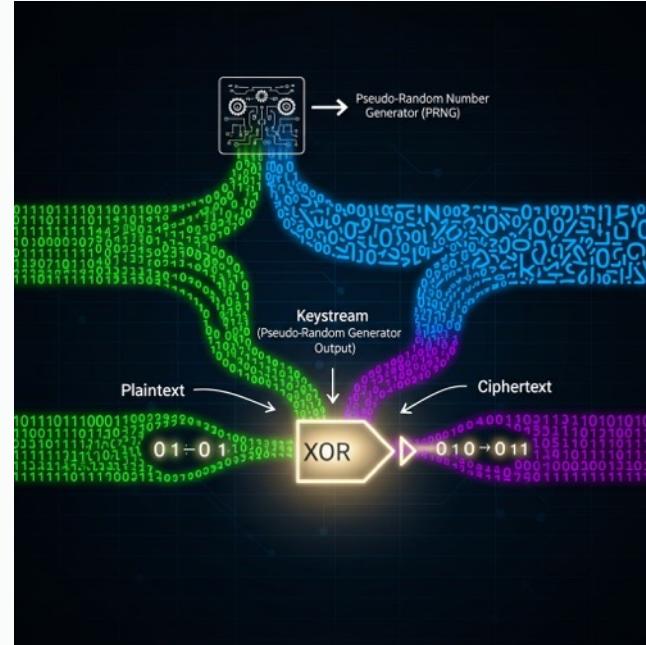
# RC-4 keystream generation

Input: permutation S of 0, 1, ..., 255

1.  $i = 0, j = 0$
2. `while (true)`
  1.  $i = (i + 1) \bmod 256$
  2.  $j = (j + S_i) \bmod 256$
  3. exchange  $S_i$  and  $S_j$
  4.  $t = (S_i + S_j) \bmod 256$
  5.  $k = S_t$  // compute XOR

at every iteration compute the XOR between k and next byte of plaintext (or ciphertext)

# Perfect ciphers and One-Time Pad



# Perfect ciphers

Plaintext space =  $\{0, 1\}^n$ ,  $D$  known

Given a ciphertext  $C$  the probability that exists  $k_2$  such that  $D_{k_2}(C) = P$  for any plaintext  $P$  is equal to the apriori probability that  $P$  is the plaintext

In other words: *the ciphertext does not reveal any information on the plaintext*

$\Pr[\text{plaintext} = P \mid \text{ciphertext} = C] = \Pr[\text{plaintext} = P]$

in short:  
 $\Pr[P|C] = \Pr[P]$

Probabilities are over the key space and the plaintext space

# Conditional probabilities

$$\Pr[P|C] = \Pr[P \wedge C] / \Pr[C]$$

(def. of cond. pr.)

$$\Pr[P|C] * \Pr[C] =$$
$$\Pr[C|P] * \Pr[P]$$

(Th. Bayes)

In a perfect cipher

$$\Pr[P|C] = \Pr[P]$$

by replacing into Bayes

$$\Pr[P] * \Pr[C] = \Pr[C|P] * \Pr[P]$$

$$\Pr[C] = \Pr[C | P]$$

# One-time pad (OTP)

AKA Vernam Cipher,  
invented in 1917 and  
patented in 1919  
while *Gilbert Vernam*  
was working at AT&T

Plaintext space:  
 $\{0,1\}^n$

Keystream space:  
 $\{0,1\}^n$

The scheme is  
symmetric;  
keystream  $K$  is  
chosen at random

$$E_K(P) = C = P \oplus K$$

$$D_K(C) = C \oplus K = P \oplus K \oplus K = P$$

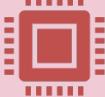
# Features of OTP



Claim: one-time pad is a perfect cipher. [given a  $k$  bit cipher text every  $k$ -bit plain text has got same probability if key is random]



Problem: size of keystream space, as show by the following



**Theorem** (Shannon): A cipher cannot be perfect if the size of its key space is less than the size of its message space

# Proof of Shannon's theorem

By contradiction

Assume #keys ( $l$ ) < #messages ( $n$ ) and consider ciphertext  $C_0$  s.t.  $\Pr[C_0] > 0$  ( $C_0$  must exist!)

For some key  $K$ , consider  $P = D_K(C_0)$ . There exist at most  $l$  (#keys) such messages (one per each key)

Choose message  $P_0$  s.t. it is not of the form  $D_K(C_0)$  (there exist  $n-l$  such messages)

Hence  $\Pr[C_0|P_0] = 0$

But in a perfect cipher  $\Pr[C_0|P_0] = \Pr[C_0] > 0$ . Contradiction

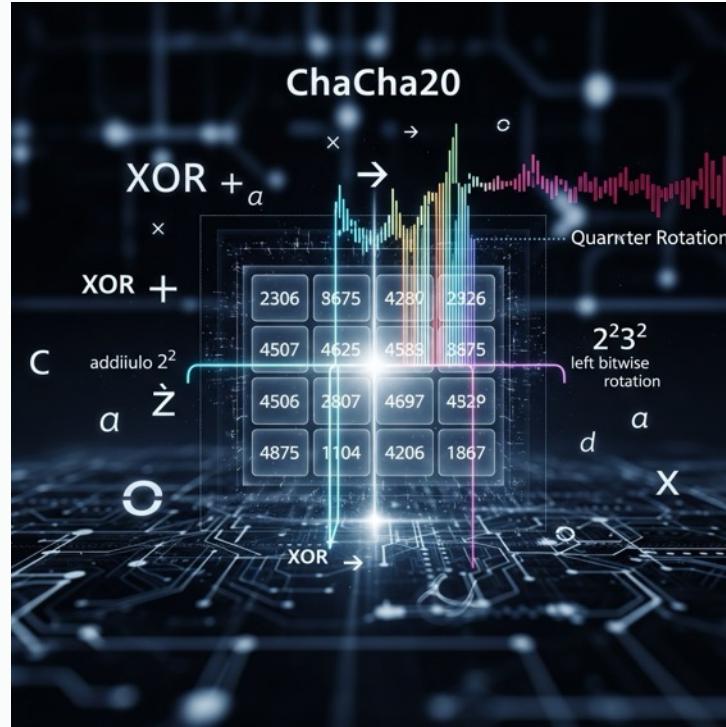
# Consequences

- OTP is perfect if
  - keystream truly random
  - keystream space not less message space
- if seed (key) is reused we get same keystream  $K$  for two different encryptions. Heavy consequences
  - $C_1 = P_1 \oplus K, C_2 = P_2 \oplus K$
  - Assume pair plaintext, ciphertext  $(P_1, C_1)$  known (KPA)
  - Attacker computes  $C_1 \oplus C_2 = P_1 \oplus K \oplus P_2 \oplus K$
  - $\oplus$  commutative and associative, so:  $C_1 \oplus C_2 = P_1 \oplus P_2 \oplus K \oplus K = P_1 \oplus P_2$  (by properties of  $\oplus$ )
  - $P_1 \oplus P_2 \oplus P_1$  (known) =  $P_2$  (Success!)

# Conclusion

- OTP is not obsolete
- it needs a truly random long keystream
  - PRNG not enough
- not to be reused
- This makes OTP unpractical
- Otherwise... OTP is theoretically unbreakable (the only with proved perfect security)

# ChaCha20



# ChaCha20, a synchronous stream cipher

## Objectives

- Understand the design and purpose of ChaCha20
- Compare it to older stream cipher models
- Learn how ChaCha20 builds on the one-time pad concept

# Background & history

- Designed in 2008 by Daniel J. Bernstein as an improvement of Salsa20
- Standardized in RFC 8439 by the IETF
- Focus: high speed, simplicity, and resistance to timing attacks
- Adopted in security-critical applications like TLS 1.3, OpenSSH, WireGuard, Signal, etc.

# Stream ciphers recap

- Stream ciphers generate a keystream and encrypt data by XORing it with plaintext
- Fast, low-latency encryption ideal for real-time applications
- Example: the One-Time Pad — perfectly secure but impractical

# From OTP to ChaCha20

- ChaCha20 mimics the One-Time Pad by generating a pseudo-random keystream
- Key idea: instead of a long random keystream, we use a small key + counter + nonce (random number to be used once) to generate keystream blocks
- XOR remains the central operation, just like OTP

# Inputs to ChaCha20

- **256-bit key:** shared secret
- **96-bit nonce:** unique per message/session
- **32-bit block counter:** changes per block to ensure keystream uniqueness
  - a block is a chunk of 512 bits of plaintext
- These inputs determine the internal state that drives the keystream

# Block structure

- 512-bit state divided into 16 32-bit words
  - 384 bits are variable and 128 are coming from constants
- Layout
  - 4 constant words
  - 8 key words
  - 1 counter word
  - 3 nonce words
- This state is transformed to produce 64-byte keystream blocks

# Quarter-round function

- Core operation: processes 4 words (from state) with **Add, Rotate, XOR (ARX)**
- Lightweight and fast on all CPUs
- Provides diffusion and non-linearity
- Designed for simplicity and timing safety

# 10 Double-rounds and output

- ChaCha20 performs 20 operations, organized as 10 double-rounds
- Each double-round consists of
  - 1 column-wise operation: applies 4 Quarter-Round functions to the columns of the  $4 \times 4$  word matrix
  - 1 diagonal-wise operation: applies 4 Quarter-Round functions to the diagonals of the matrix
- 1 double-round = 8 Quarter-Round functions
- These transformations ensure that all 16 words of the state are thoroughly mixed in a defined pattern
- After all 10 double-rounds: add the final state to the original state ( $\text{mod } 2^{32}$ ) to get 64 bytes of keystream

# Encryption with XOR

- **ciphertext = plaintext  $\oplus$  keystream**
- Same operation used for decryption: **plaintext = ciphertext  $\oplus$  keystream**
- Stateless per block: no chaining between blocks required

# Security goals of ChaCha20

- Designed to resist
  - Key recovery
  - Keystream prediction
  - Timing attacks
- No known practical attacks on full 10 double-round version
- High performance across platforms, even without hardware support

# Real-world usage

- TLS 1.3: used as an option (in addition to others) for secure web traffic
- Signal, WhatsApp: encrypt messages using ChaCha20 (or other, depending on context)
- WireGuard VPN: lightweight and fast encryption
- OpenSSH: option for encrypted shell sessions, now less used