



# Vector instructions: SIMD

# Single Instruction Multiple Data




The SIMD instruction executes 8 DP floating point operations , corresponding to a whole cache line

# Vectorization trend



SIMD instruction moved from 64 bits of MMX to 512 bits of AVX512

# Vector size



A single SIMD register can be used with different data types

# Vector register file

	511	256	255	128	127	0
ZMM0		YMM0	XMM0			
ZMM1		YMM1	XMM1			
ZMM2		YMM2	XMM2			
ZMM3		YMM3	XMM3			
ZMM4		YMM4	XMM4			
ZMM5		YMM5	XMM5			
ZMM6		YMM6	XMM6			
ZMM7		YMM7	XMM7			
ZMM8		YMM8	XMM8			
ZMM9		YMM9	XMM9			
ZMM10		YMM10	XMM10			
ZMM11		YMM11	XMM11			
ZMM12		YMM12	XMM12			
ZMM13		YMM13	XMM13			
ZMM14		YMM14	XMM14			
ZMM15		YMM15	XMM15			
ZMM16		YMM16	XMM16			
ZMM17		YMM17	XMM17			
ZMM18		YMM18	XMM18			
ZMM19		YMM19	XMM19			
ZMM20		YMM20	XMM20			
ZMM21		YMM21	XMM21			
ZMM22		YMM22	XMM22			
ZMM23		YMM23	XMM23			
ZMM24		YMM24	XMM24			
ZMM25		YMM25	XMM25			
ZMM26		YMM26	XMM26			
ZMM27		YMM27	XMM27			
ZMM28		YMM28	XMM28			
ZMM29		YMM29	XMM29			
ZMM30		YMM30	XMM30			
ZMM31		YMM31	XMM31			

AVX512 registers (ZMM[0-31]) are a superset of AVX/AVX2 YMM[0-15] registers, and of SSE XMM[0-15] registers

# check for SIMD capabilities

```
[v] lscpu | grep 'sse\|avx'  
Flags:  
fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx pdp  
d aperf mperf tsc_known_freq pni pclmulqdq dtes64 monitor ds_cpl vmx smx est tm2 sse3 sdbg fma cx16 xtpr pdcm pcid dca sse4_1 sse4_2 x2apic movbe popcnt tsc_deadline_timer  
_ppin cdp_l2 ssbd mba ibrs ibpb stibp ibrs_enhanced tpr_shadow flexpriority ept vpid ept_ad fsgsbase tsc_adjust bmi1 avx2 smep bmi2 erms invpcid cqmq rdt_a avx512f avx512dq  
xsavexc xgetbv1 xsaves cqmq_llc cqmq_occup_llc cqmq_mbm_total cqmq_mbm_local split_lock_detect user_shstk avx_vnni avx512_bf16 wbnoinvd dtherm ida arat pln pts hwp hwp_act_wind
```

# Compiler flags

gcc/clang compilers can exploit vectorization instructions if the right flags are used

```
-msse  
-msse2  
...  
-msse4.2  
-mavx  
-mavx2  
-mavx512f
```

A full list of x86 instructions is available here: <https://www.intel.com/content/www/us/en/docs/intrinsics-guide/index.html>

# vector intrinsics types

the <immintrin.h> header provides intrinsinc instructions and types

INTEGER types					
	types name	int 8bits	int 16bits	int 32bits	int 64bits
64bits	<code>_m64</code>	8x	4x	2x	1x
128bits	<code>_m128i</code>	16x	8x	4x	2x
256bits	<code>_m256i</code>	32x	16x	8x	4x
512bits	<code>_m512i</code>	64x	32x	16x	8x

FLOATING-POINT types			
	types name	single precision	double precision
128bits	<code>_m128</code>	4x	-
	<code>_m128d</code>	-	2x
256bits	<code>_m256</code>	8x	-
	<code>_m256d</code>	-	4x
512bits	<code>_m512</code>	16x	-
	<code>_m512d</code>	-	8x

# Intrinsic Name convention

The naming of SIMD instructions follows a pretty systematic set of rules:

```
_mm[bits]_[operation]_[type][modifier]
```

## 1. Prefix (\_mm, \_mm256, \_mm512)

\_mm → 128-bit (xmm register)

\_mm256 → 256-bit (ymm register)

\_mm512 → 512-bit (zmm register)

## 2. Operation

add, sub, mul, div, and, or, xor, etc.

Sometimes longer: fma (fused multiply-add), cvtepi32\_ps  
(convert int32 to float), etc.

# Intrinsic Name convention

## 3. Type Suffix

- **ps** = packed single-precision float (float32)
- **pd** = packed double-precision float (float64)
- **epi8/16/32/64** = signed integers of given width
- **epu8/16/32/64** = unsigned integers of given width

## 4. Optional Modifiers

- **s** = scalar (lowest element only)
- **u** = unaligned load/store (e.g., \_mm\_loadu\_si128)

# Intrinsic Name examples

1. `_mm256_add_ps`

`_mm256` → 256-bit vector (ymm)

`add` → addition

`ps` → packed single-precision floats

adds 8 floats in parallel.

2. `_mm512_cvtepi32_pd`

`_mm512` → 512-bit (zmm)

`cvt` → convert epi32 → from signed 32-bit integers

`pd` → to double-precision floats

converts 8 int32 → 8 float64.

# example

```
int main() {
    float a[N], b[N], c[N];
    // Initialize arrays
    for (int i = 0; i < N; i++) {
        a[i] = i * 1.0f;  b[i] = (N - i) * 1.0f;
    }
    // Perform vectorized addition using AVX
    for (int i = 0; i < N; i += 8) {
        // Load 8 floats from each array
        __m256 va = _mm256_load_ps(&a[i]);
        __m256 vb = _mm256_load_ps(&b[i]);
        // Add the two vectors
        __m256 vc = _mm256_add_ps(va, vb);
        // Store result into c
        _mm256_store_ps(&c[i], vc);
    }
}
```

# GCC/Clang Extended Vector Types

There are many different ways of using vector operations in C++ code.

Many programmers use intrinsic instructions that provide very low-level access to the vector operations.

However, there is also a much more programmer-friendly approach: many modern C++ compilers support convenient **extended vector types**.

# GCC/Clang Vector Types

This typedef let you write arithmetic on vector registers directly:

```
typedef float float8_t __attribute__ ((vector_size (8 * sizeof(float))));
```

The following lines are compiled to AVX instructions:

```
float8_t a, b, c;  
c = a + b; // compiler emits one SIMD add (e.g. AVX vaddps)
```

# Warning: proper memory alignment needed

If we do dynamic memory allocation of `float8_t`, we **must make sure** that it points to a block of memory that is properly aligned.

Solutions:

- use `posix_memalign()` function;
- with C11 you can use `aligned_alloc()`
- with C17++ you can use

```
// compile with -std=c++17
std::vector<float8_t> x(10);
```

# Code vectorization

# Code vectorization

- The software changes required to exploit SIMD instructions are known as **code vectorization**.
- Initially, SIMD instructions were programmed in assembly. Later, special compiler intrinsics, which are small functions providing a one-to-one mapping to SIMD instructions, were introduced.
- ***Autovectorization***: today all the major compilers support autovectorization for the popular processors, i.e., they can generate SIMD instructions straight from high-level code written in C/C++, Java, Rust, and other languages.

# Code vectorization

However, in some cases, autovectorization does not succeed without intervention by a software engineer

- loop indices may *overflow*, preventing certain loop transformations.
- *pointers* in the program may overlap
- processors that don't have efficient vector instructions for operations . (e.g. predicated operations) are not available on most processors.

# Code vectorization

- Modern compilers can report whether a certain loop was vectorized. If the compiler cannot vectorize a loop, it is also able to tell the reason why it failed.
- Sometimes you can use compiler hints to drive vectorization
- When you cannot rely on compiler autovectorization or compiler hints, code can instead be written using *compiler intrinsics*. In most cases, compiler intrinsics provide a 1-to-1 mapping to assembly instructions.

# Code vectorization

The vectorizer is usually structured in three phases:

- **Legality-check:** in this phase, the compiler checks if it is legal to transform the loop using vectors.
  - checks that the iterations of the loop are consecutive,
  - ensures that all memory and arithmetic operations can be widened into consecutive operations.
  - check that the order of operations will be preserved.

# Code vectorization

**Profitability-check:** the vectorizer checks if a transformation is profitable.

- It compares different vectorization widths and figures out which one would be the fastest to execute.
- The vectorizer uses a cost model to predict the cost of different operations, such as scalar add or vector load.
- It takes into account the added instructions that shuffle data into registers, and estimate the cost of the loop guards that ensure that preconditions that allow vectorizations are met.

# Code vectorization

**Transformation:** The vectorizer actually transforms the code

- This process also includes the insertion of guards that enable vectorization. If the loop use an unknown iteration count, the compiler has to generate a scalar version of the loop (remainder), in addition to the vectorized version of the loop, to handle the last few iterations.
- The compiler also has to check if pointers don't overlap, etc.

All of these transformations are done using information that is collected during the legality check phase.

# Discovering Vectorization Opportunities

A vectorization opportunity should start by analyzing hot loops in the program and checking why vectorization was not applied.

A good indicator for potential vectorization opportunities is a high **retiring rate** (above 80%).

Perhaps a workload executes a lot of simple instructions that can be replaced by vector instructions. In such situations, high retiring metric doesn't translate into high performance.

# Example 1: illegal vectorization

**Listing 9.13** Vectorization: read-after-write dependence.

```
void vectorDependence(int *A, int n) {
    for (int i = 1; i < n; i++)
        A[i] = A[i-1] * 2;
}
```

This code cannot be vectorized due to read-after-write

```
$ clang -Rpass-analysis=loop-vectorize -march=native -O3 vectorDependence.c
```

```
Backward loop carried data dependence. Memory location is the same as accessed at vectorDependence.c:15:24 [-Rpass-analysis=loop-vectorize]
15 |             A[i] = A[i-1] * 2;
```

# Example 1: illegal vectorization

```
void vectorDependence_unrolled(int *A) {  
  
    // First element is assumed initialized  
    for (int i = 1; i+7 < N; i += 8) {  
        A[i]    = A[i-1] * 2;  
        A[i+1]  = A[i-1] * 4;  
        A[i+2]  = A[i-1] * 8;  
        A[i+3]  = A[i-1] * 16;  
        A[i+4]  = A[i-1] * 32;  
        A[i+5]  = A[i-1] * 64;  
        A[i+6]  = A[i-1] * 128;  
        A[i+7]  = A[i-1] * 256;  
    }  
}
```

does not work  
(openMP loop)

```
$ clang -Rpass-analysis=loop-vectorize -march=native -O3 vectorDependence.c
```

```
Backward loop carried data dependence. Memory location is the same as accessed at vectorDependence.c:27:18 [-Rpass-analysis=loop-vectorize]  
34 |         A[i+7] = A[i-1] * 256;
```

# Example 1: illegal vectorization

```
void vectorDependence_unrolled_mm(int *A) {  
  
    __m256i base = _mm256_set1_epi32(A[0]); // broadcast A[0] to 8 ints  
    const int power2_const[8] = {2, 4, 8, 16, 32, 64, 128, 256};  
    __m256i power2 = _mm256_loadu_si256((__m256i*)power2_const);  
  
    __m256i vals = _mm256_mullo_epi32(base, power2);  
    _mm256_storeu_si256((__m256i*)&A[1], vals);  
    base = _mm256_set1_epi32(256); // broadcast 256 to 8 ints  
    for (int i = 9; i < N; i += 8) {  
        // Compute the values by multiplying with A[i-8]  
        vals = _mm256_mullo_epi32(vals, base);  
        _mm256_storeu_si256((__m256i*)&A[i], vals);  
    }  
}
```

(dif

tructions

```
$ clang -Rpass-analysis=loop-vectorize -march=native -O3 vectorDependence.c
```

```
vectorDependence.c:49:5: remark: loop not vectorized: loop control flow is not understood by vectorizer [-Rpass-analysis=loop-vectorize]  
49 |     for (int i = 9; i < N; i += 8) {
```

# Example 1: illegal vectorization

```
void vectorDependence_unrolled_mm(int *A) {  
  
    __m256i base = _mm256_set1_epi32(A[0]); // broadcast A[0] to 8 ints  
    const int power2_const[8] = {2, 4, 8, 16, 32, 64, 128, 256};  
    __m256i power2 = _mm256_loadu_si256((__m256i*)power2_const);  
  
    __m256i vals = _mm256_mullo_epi32(base, power2);  
    _mm256_storeu_si256((__m256i*)&A[1], vals);  
    base = _mm256_set1_epi32(256); // broadcast 256 to 8 ints  
    for (int i = 9; i < N; i += 8) {  
        // Compute the values by multiplying with A[i-8]  
        vals = _mm256_mullo_epi32(vals, base);  
        _mm256_storeu_si256((__m256i*)&A[i], vals);  
    }  
}
```

(dif

tructions

```
$ clang -Rpass-analysis=loop-vectorize -march=native -O3 vectorDependence.c
```

```
vectorDependence.c:49:5: remark: loop not vectorized: loop control flow is not understood by vectorizer [-Rpass-analysis=loop-vectorize]  
49 |     for (int i = 9; i < N; i += 8) {
```

Why?

# Example 1: illegal vectorization

```
void vectorDependence_vectorizable(int *A) {
    int base = A[0];
    for (int i = 1; i < N; i++) {
        A[i] = base << i; // equivalent to A[0] * (2^i)
    }
}
```

algorithm  
(IP loop)

```
$ clang -Rpass-analysis=loop-vectorize -march=native -O3 vectorDependence.c
```

```
vectorDependence.c:59:5: remark: vectorized loop (vectorization width: 8, interleaved count: 4) [-Rpass=loop-vectorize]
```

# Example 1: illegal vectorization

```
void vectorDependence_vectorizable(int *A) {
    int base = A[0];
    for (int i = 1; i < N; i++) {
        A[i] = base << i; // equivalent to A[0] * (2^i)
    }
}
```

hm  
loop)

```
$ clang -Rpass-analysis=loop-vectorize -march=native -O3 vectorDependence.c
```

```
'vectorDependence.c:59:5: remark: vectorized loop (vectorization width: 8, interleaved count: 4) [-Rpass=loop-vectorize]
```

width:8 ☐ 8 integer in the same SIMD

interleaving:4 ☐ 4 SIMD instructions for each loop iteration

# Example 1: illegal vectorization

## Assembly of vectorized loop

The screenshot shows the Compiler Explorer interface with two panes. The left pane displays the C source code, and the right pane shows the generated assembly code.

**C source #1:**

```
1 #include <stdio.h>
2 #include <stdint.h>
3
4 # define N 1025
5
6 void vectorDependence_vectorizable(int *A) {
7     int base = A[0];
8     for (int i = 1; i < N; i++) {
9         A[i] = base << i; // equivalent to A[0] * (2^i)
10    }
11 }
12
13
14 int main() {
15     int A[N];
16
17     // Initialize first element
18     A[0] = 1;
19
20     vectorDependence_vectorizable(A);
21
22     // Print results
23     printf("Resulting array:\n");
24     //for (int i = 0; i < N; i++) {
25     //    printf("%d ", A[N-1]);
26     //}
27     printf("\n");
28
29     return 0;
30 }
```

**x86-64 clang 20.1.0 (Editor #1)**

```
-O3 -mavx2
```

```
17 .byte 8
18 vectorDependence_vectorizable:
19 vpbroadcastd ymm0, dword ptr [rdi]
20 vpmovsxbd ymm1, qword ptr [rip + .LCPI0_5]
21 xor eax, eax
22 vpbroadcastd ymm2, dword ptr [rip + .LCPI0_1]
23 vpbroadcastd ymm3, dword ptr [rip + .LCPI0_2]
24 vpbroadcastd ymm4, dword ptr [rip + .LCPI0_3]
25 vpbroadcastd ymm5, dword ptr [rip + .LCPI0_4]
26 .LBB0_1:
27 vpaddd ymm6, ymm1, ymm2
28 vpaddd ymm7, ymm1, ymm3
29 vpaddd ymm8, ymm1, ymm4
30 vpslld ymm9, ymm0, ymm1
31 vpslld ymm6, ymm0, ymm6
32 vpslld ymm7, ymm0, ymm7
33 vpslld ymm8, ymm0, ymm8
34 vmovdqu ymmword ptr [rdi + 4*rax + 4], ymm9
35 vmovdqu ymmword ptr [rdi + 4*rax + 36], ymm6
36 vmovdqu ymmword ptr [rdi + 4*rax + 68], ymm7
37 vmovdqu ymmword ptr [rdi + 4*rax + 100], ymm8
38 add rax, 32
39 vpaddd ymm1, ymm1, ymm5
40 cmp rax, 1024
41 jne .LBB0_1
42 vzeroupper
43 ret
44
45 .LCPI1_1:
```

# Example 2: floating-point arithmetic

**Listing 9.14** Vectorization: floating-point arithmetic.

```
1 // a.cpp
2 float calcSum(float* a, unsigned N) {
3     float sum = 0.0f;
4     for (unsigned i = 0; i < N; i++) {
5         sum += a[i];
6     }
7     return sum;
8 }
```

This code cannot be vectorized due to not associative FP

O

```
$ clang -Rpass-analysis=loop-vectorize -march=native -O3 calcsum.c
```

```
calcsum.c:12:13: remark: loop not vectorized: cannot prove it is safe to reorder floating-point operations; allow reordering by specifying
#pragma clang loop vectorize(enable)' before the loop or by providing the compiler option '-ffast-math' [-Rpass-analysis=loop-vectorize]
```

# Example 2: floating-point arithmetic

**Listing 9.14** Vectorization: floating-point arithmetic.

```
1 // a.cpp
2 float calcSum(float* a, unsigned N) {
3     float sum = 0.0f;
4     for (unsigned i = 0; i < N; i++) {
5         sum += a[i];
6     }
7     return sum;
8 }
```

This code cannot be vectorized due to not associative FP

```
$ clang -Rpass=loop-vectorize -Rpass-analysis=loop-vectorize \
        -ffast-math -march=native -O3 calcsum.c
```

```
calcsum.c:11:5: remark: vectorized loop (vectorization width: 8, interleaved count: 4) [-Rpass=loop-vectorize]
11 |     for (unsigned i = 0; i < N; i++) {
|     ^
calcsum.c:21:5: remark: vectorized loop (vectorization width: 8, interleaved count: 2) [-Rpass=loop-vectorize]
21 |     for (unsigned i = 0; i < N; i++) {
|     ^
calcsum.c:11:5: remark: vectorized loop (vectorization width: 8, interleaved count: 4) [-Rpass=loop-vectorize]
11 |     for (unsigned i = 0; i < N; i++) {
|     ^
```

# Example 3: memory overlap

---

**Listing 9.15** a.c

```
1 void foo(float* a, float* b, float* c, unsigned N) {  
2     for (unsigned i = 1; i < N; i++) {  
3         c[i] = b[i];  
4         a[i] = c[i-1];  
5     }  
6 }
```

---

When compilers cannot prove that a loop operates on arrays with non-overlapping memory regions, they usually choose to be on the safe side

```
$ clang -Rpass=loop-vectorize -Rpass-analysis=loop-vectorize \  
      -march=native -O3 overlap.c
```

# Example 3: memory overlap

---

**Listing 9.15** a.c

---

```
1 void foo(float* a, float* b, float* c, unsigned N) {  
2     for (unsigned i = 1; i < N; i++) {  
3         c[i] = b[i];  
4         a[i] = c[i-1];  
5     }  
6 }
```

---

When compilers cannot prove that a loop operates on arrays with non-overlapping memory regions, they usually choose to be on the safe side

```
$ clang -Rpass=loop-vectorize -Rpass-analysis=loop-vectorize \  
      -march=native -O3 overlap.c
```

```
overlap.c:8:16: remark: loop not vectorized: unsafe dependent memory operations in loop. Use #pragma clang loop distribute(enable) to allow  
loop distribution to attempt to isolate the offending operations into a separate loop  
Forward loop carried data dependence that prevents store-to-load forwarding. Memory location is the same as accessed at overlap.c:7:9 [-Rpa  
ss-analysis=loop-vectorize]  
8 |         a[i] = c[i-1];
```

# Example 3: memory overlap

---

**Listing 5.7** a.c

```
1 void foo(float* __restrict__ a,
2         float* __restrict__ b,
3         float* __restrict__ c,
4         unsigned N) {
5     for (unsigned i = 1; i < N; i++) {
6         c[i] = b[i];
7         a[i] = c[i-1];
8     }
9 }
```

When a developer knows arrays a, b, and c do not overlap, it can insert the \_\_restrict\_\_ keyword

```
$ clang -Rpass=loop-vectorize -Rpass-analysis=loop-vectorize \
        -march=native -O3 overlap.c
```

# Example 3: memory overlap

**Listing 5.7** a.c

```
1 void foo(float* __restrict__ a,
2          float* __restrict__ b,
3          float* __restrict__ c,
4          unsigned N) {
5     for (unsigned i = 1; i < N; i++) {
6         c[i] = b[i];
7         a[i] = c[i-1];
8     }
9 }
```

When a developer knows arrays a, b, and c do not overlap, it can insert the `__restrict__` keyword

```
$ clang -Rpass=loop-vectorize -Rpass-analysis=loop-vectorize \
        -march=native -O3 overlap.c
```

```
overlap.c:6:5: remark: vectorized loop (vectorization width: 8, interleaved count: 4) [-Rpass=loop-vectorize]
6 |     for (unsigned i = 1; i < N; i++) {
```

# Example 4: Vectorization Is Not Beneficial

---

**Listing 9.16** Vectorization: not beneficial.

```
1 // a.cpp
2 void stridedLoads(int *A, int *B, int n) {
3     for (int i = 0; i < n; i++)
4         A[i] += B[i * 3];
5 }
```

In some cases, the compiler can vectorize the loop but decide that doing so is not profitable

```
$ clang -c -O3 -mavx strideload.c -Rpass-missed=loop-vectorize
```

```
strideload.c:7:5: remark: the cost-model indicates that vectorization is not beneficial [-Rpass-missed=loop-vectorize]
 7 |     for (int i = 0; i < n; i++)
  |     ^
```

# Programming style for better vectorization

- Use `__restrict__` attribute on pointers
- Use pragmas carefully:
  - to inform compiler
  - `#pragma unroll` can prevent loop vectorization
- avoid exceptions and break
- Use contiguous and aligned memory access (e.g. use SoA)
- Define local variables directly in the loop
- Avoid function call (use inlining)

# Compiler flags

- ftree-vectorize: enables automatic vectorization of loops
- march=native: generate code optimized for the CPU of the machine you are compiling on
- fopt-info-vec: generate vectorization info.
- fopt-info-vec-missed : reports loops that couldn't be vectorized.