

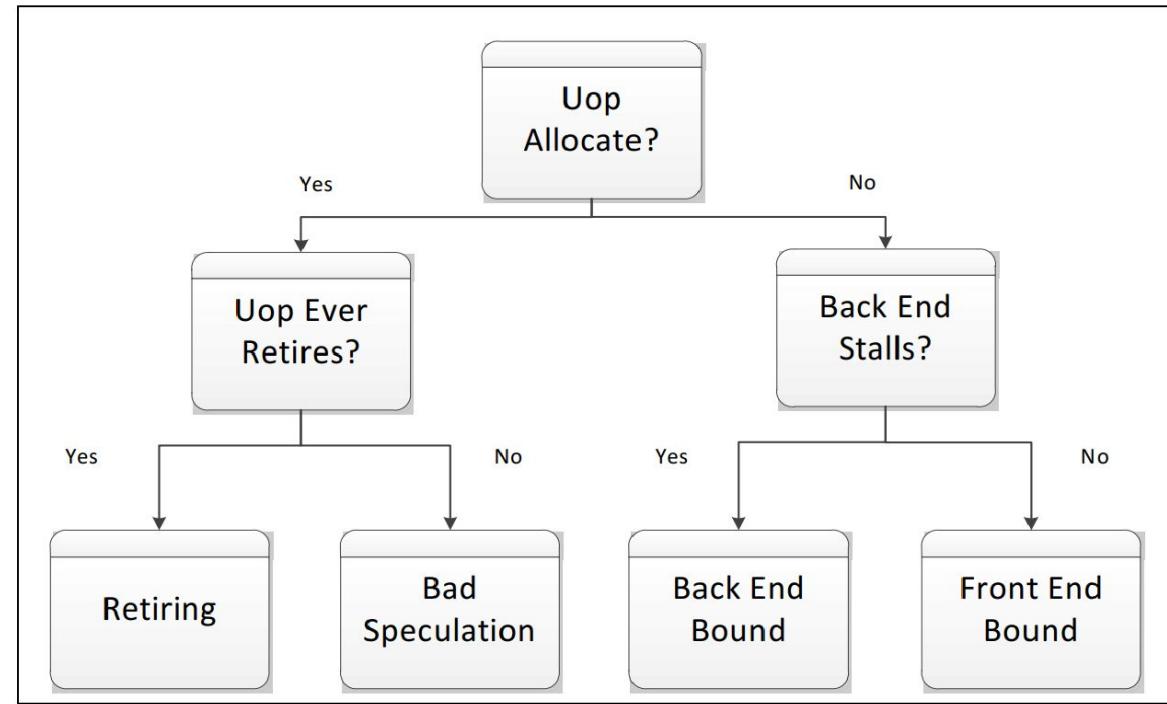
Top-down Microarchitecture Analysis

Top-down Microarchitecture Analysis

- Top-down Microarchitecture Analysis (TMA) methodology identifies ineffective usage of CPU microarchitecture by a program.
- It characterizes the **bottleneck** of a workload and allows locating the **exact place** in the source code where it occurs.
- It abstracts away the intricacies of the CPU microarchitecture and is relatively easy to use even for inexperienced developers.

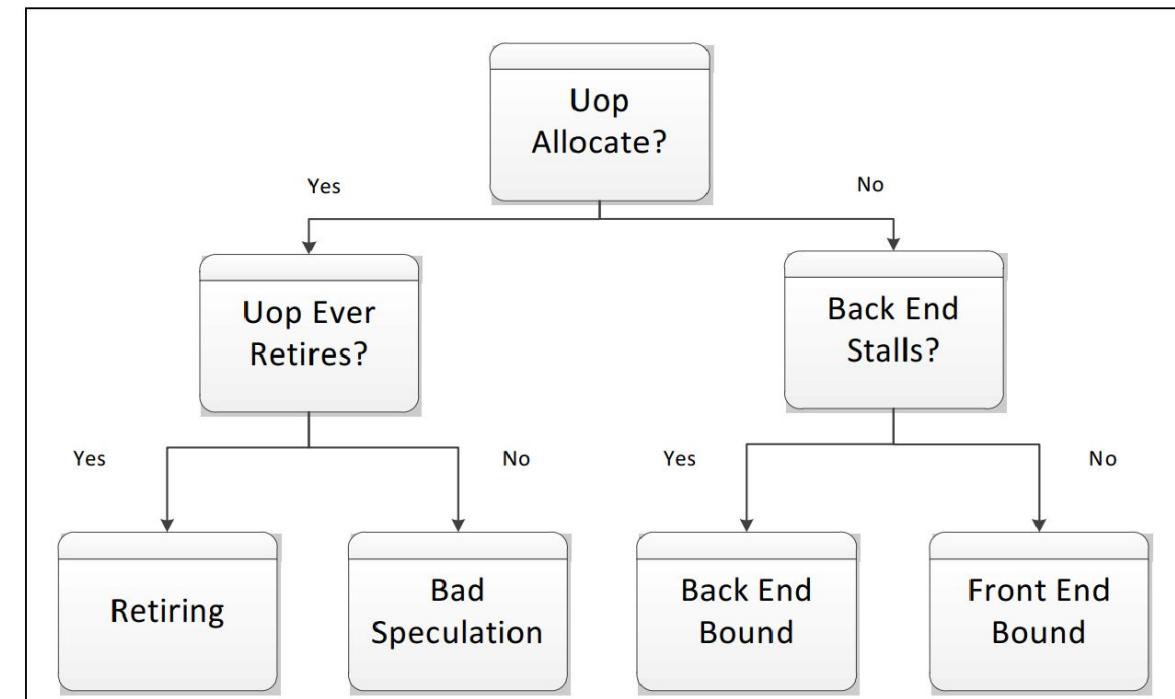
Top-down Microarchitecture Analysis

- TMA is a **hierarchical** organization of event-based metrics.
- It identifies the dominant performance bottlenecks in an application.
- It shows, on average, how well the CPU's pipelines are being utilized while running an application.



Top-down Microarchitecture Analysis

- The allocation of a μ op can fail for one of two reasons:
 - The CPU were not able to fetch and decode it (**Frontend Bound**),
 - the **Backend** was overloaded with work, (**Backend Bound**).
- If a μ op was allocated but never retired, this means it came from a mispredicted path (**Bad Speculation**).
- Finally, **Retiring** represents a normal execution.

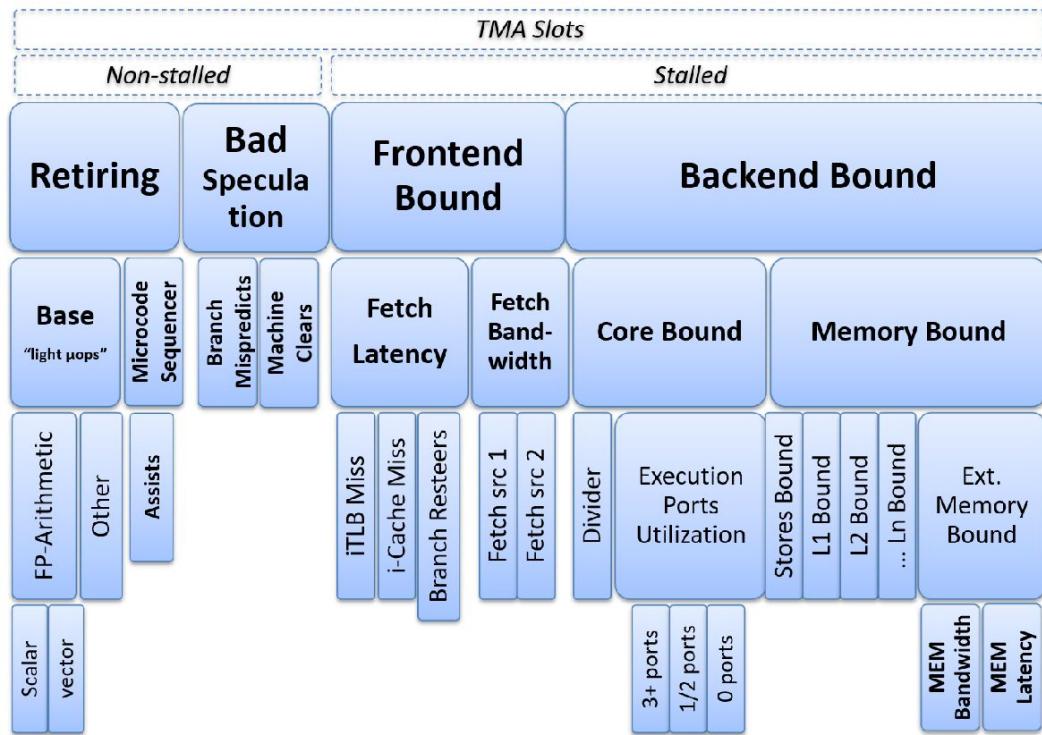


TMA hierarchy

TMA monitors a specific set of events and calculates metrics.

Using these metrics, TMA characterizes the program by assigning it to one of the four high-level buckets.

Each of the four high-level categories has several nested levels, which *CPU vendors* may choose to implement differently.



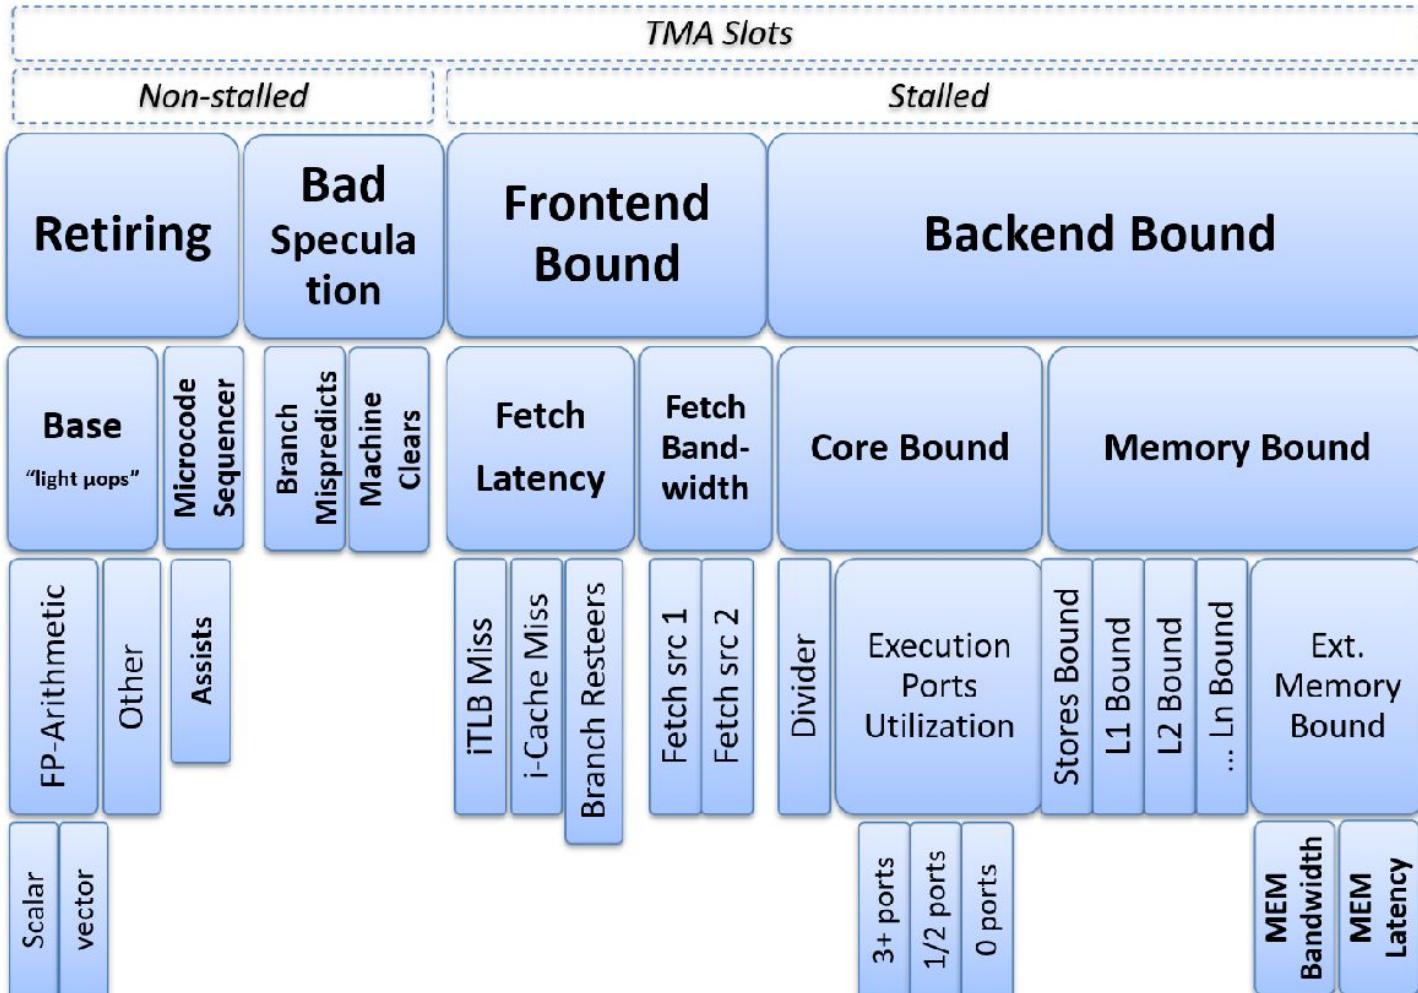
TMA hierarchy

Front End Bound

- Code Duplication
- Code Layout (Locality)
- Frequent Branching
- Unnecessary Work

Back End Bound

- *Core Bound*
 - Data Dependencies
 - Divisions and Special Functions
- *Memory Bound*
 - False Sharing
 - Scattered Memory Accesses



TMA workflow

- The workflow is designed to “drill down” to lower levels of the TMA hierarchy until we get to the very specific classification of the bottleneck.
 - For example, if a big portion of the program execution was stalled by memory accesses (**Backend Bound**) we can collect metrics specific to the Memory Bound bucket only, identifying an L3 Bound.
- The top two levels of TMA metrics are expressed in the percentage of all pipeline slots. It allows TMA to give an accurate representation of CPU microarchitecture utilization.
 - starting from Level 3, buckets may be expressed in clocks and stalls cycles (we don’t need to compare with other TMA buckets)

TMA workflow

TMA workflow is split in three stages:

1. **Bottleneck identification:** we collect events to classify the bottleneck.
2. **Bottleneck localization:** we sample on the event that characterize the bottleneck to identify the line in the source code that contributes to the bottleneck.
3. **Bottleneck removal:** we change the code to remove the bottleneck (if possible)
4. We go back to 1.

Case Study: Reduce The Number of Cache Misses with TMA

- We used a simple benchmark to show the TMA workflow. Since TMA is very effective even if you see the application for the first time we don't show the source code⁽¹⁾.
- Here is a short description: the benchmark allocates a 200 MB array on the heap, then enters a loop of 100M iterations. On every iteration of the loop, it generates a random index into the allocated array, performs some dummy work, and then reads the value from that index.
- To obtain TMA metrics Level 2, 3, and further, we will use the toplev tool that is a part of pmu-tools⁽²⁾

(1) Code is here: https://github.com/dendibakh/dendibakh.github.io/tree/master/_posts/code/TMAM

(2) PMU tools - <https://github.com/andikleen/pmu-tools>

Step 1: Identify the Bottleneck

As a first step, we run our microbenchmark and collect a limited set of events that will help us calculate Level 1 metrics.

```
$ perf stat -- ./benchmark.exe
...
TopdownL1  (cpu_core) # 53.4 % tma_backend_bound <=>
# 0.2 % tma_bad_speculation
# 13.8 % tma_frontend_bound
# 32.5 % tma_retiring
...
```

Thus, the application is bound by the **CPU backend**

Step 1: Identify the Bottleneck

Now we collect L2 metrics.

```
$~/pmu-tools/toplev.py ....  
...  
# Level 1  
S0-C0 Frontend_Bound:           13.92 % Slots  
S0-C0 Bad_Speculation:          0.23 % Slots  
S0-C0 Backend_Bound:            53.39 % Slots  
S0-C0 Retiring:                 32.49 % Slots  
# Level 2  
S0-C0 Frontend_Bound.FE_Latency: 12.11 % Slots  
S0-C0 Frontend_Bound.FE_Bandwidth: 1.84 % Slots  
S0-C0 Bad_Speculation.Branch_Mispred: 0.22 % Slots  
S0-C0 Bad_Speculation.Machine_Clears: 0.01 % Slots  
S0-C0 Backend_Bound.Memory_Bound:   44.59 % Slots <==  
S0-C0 Backend_Bound.Core_Bound:    8.80 % Slots  
...  
...
```

Step 1: Identify the Bottleneck

Now we collect L3 metrics.

```
$~/pmu-tools/toplev.py ....  
...  
# Level 3  
S0-C0-T0 BE_Bound.Mem_Bound.L1_Bound:    4.39 % Stalls  
S0-C0-T0 BE_Bound.Mem_Bound.L2_Bound:    2.42 % Stalls  
S0-C0-T0 BE_Bound.Mem_Bound.L3_Bound:    5.75 % Stalls  
S0-C0-T0 BE_Bound.Mem_Bound.DRAM_Bound:   47.11 % Stalls <==  
S0-C0-T0 BE_Bound.Mem_Bound.Store_Bound:   0.69 % Stalls  
S0-C0-T0 BE_Bound.Core_Bound.Divider:    8.56 % Clocks  
S0-C0-T0 BE_Bound.Core_Bound.Ports_Util:  11.31 % Clocks...
```

Step 1: Identify the Bottleneck

- The bottleneck is in **DRAM_Bound**: many memory accesses miss in all levels of caches and go all the way down to the main memory.
- We can confirm this collecting the absolute number of L3 cache misses for the program.

Step 1: Identify the Bottleneck

- The bottleneck is in **DRAM_Bound**: many memory accesses miss in all levels of caches and go all the way down to the main memory.
- We can confirm this collecting the absolute number of L3 cache misses for the program.

```
$ perf stat -e cycles,cycle_activity.stalls_l3_miss -- ./benchmark.exe  
32226253316 cycles  
19764641315 cycle_activity.stalls_l3_miss
```

Step 1: Identify the Bottleneck

- The bottleneck is in **DRAM_Bound**: many memory accesses miss in all levels of caches and go all the way down to the main memory.

- We can confirm this collecting data for the program.

The CYCLE_ACTIVITYSTALLS_L3_MISS event counts cycles when execution stalls, while the L3 cache miss demand load is outstanding

```
$ perf stat -e cycles,cycle_activity.stalls_l3_miss -- ./benchmark.exe  
32226253316 cycles  
19764641315 cycle_activity.stalls_l3_miss
```

Step 1: Identify the Bottleneck

- The bottleneck is in **DRAM_Bound**: many memory accesses miss in all levels of caches and go all the way down to the main memory.
- We can confirm this collecting CPU is in stall for 60% of clock cycles for the program.

```
$ perf stat -e cycles,cycle_activity.stalls_13_miss -- ./benchmark.exe
32226253316 cycles
19764641315 cycle_activity.stalls_13_miss
```

Step 2: Locate the Place in the Code

- The second step in the TMA process is to locate the place in the code sampling using the event related to the bottleneck.
- The recommended way to find such an event is to run `toplev` with the `--show-sample` option that suggests the perf record command line
- Correspondence between performance bottlenecks and performance events can be found with the TMA metrics table⁽¹⁾.

(1) TMA metrics - https://github.com/intel/perfmon/blob/main/TMA_Metrics.xlsx

Step 2: Locate the Place in the Code

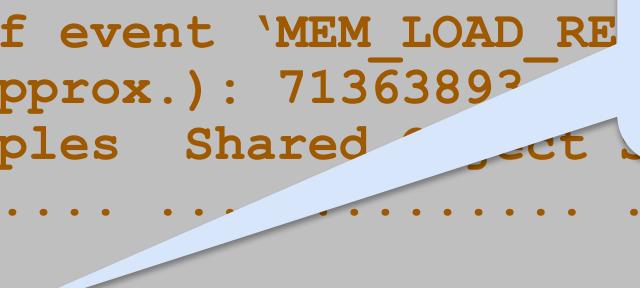
In our case, we sample on `MEM_LOAD_RETIRE.L3_MISS_PS` precise event.

```
$ perf record -e cpu/event=0xd1,umask=0x20,name=MEM_LOAD_RETIRE.L3_MISS/ooo -- ./benchmark.exe
$ perf report -n --stdio
...
# Samples: 33K of event 'MEM_LOAD_RETIRE.' L3_MISS
# Event count (approx.): 71363893
# Overhead    Samples  Shared Object Symbol
# ..... .
#
99.95% 33811  benchmark.exe  [.] foo
  0.03%    52  [kernel]      [k] get_page_from_freelist
  0.01%     3  [kernel]      [k] free_pages_prepare
  0.00%     1  [kernel]      [k] free_pcpages_bulk
```

Step 2: Locate the Place in the Code

In our case, we sample on `MEM_LOAD_RETIRE.L3_MISS_PS` precise event.

```
$ perf record -e cpu/event=0xd1,umask=0x20,name=MEM_LOAD_RETIRE.L3_MISS/ooo -- ./benchmark.exe
$ perf report -n --stdio
...
# Samples: 33K of event 'MEM_LOAD_RE
# Event count (approx.): 71363892
# Overhead    Samples  Shared Object S
# ..... . . . . . . . . . . . . . . .
#
99.95% 33811  benchmark.exe  [.] foo
  0.03%   52  [kernel]      [k] get_page_from_freelist
  0.01%    3  [kernel]      [k] free_pages_prepare
  0.00%    1  [kernel]      [k] free_pcpages_bulk
```



Almost all L3 misses are caused by memory accesses in function `foo` inside executable `benchmark.exe`.

Step 2: Locate the Place in the Code

In our case, we sample on `MEM_LOAD_RETIRE.L3_MISS_PS` precise event.

```
$ perf record -e cpu/event=0xd1,umask=0x20,name=MEM_LOAD_RETIRE.L3_MISS/ooo -- ./benchmark.exe
$ perf report -n --stdio
...
# Samples: 33K of event 'MEM_LOAD_RETIRE.' L3_MISS
# Event count (approx.): 71363893
# Overhead    Samples  Shared Object Symbol
# ..... . . . . . . . . . . . . . . .
#
99.95% 33811  benchmark.exe  [.] foo
  0.03%   52  [kernel]      [k] get_page_from_freelist
  0.01%    3  [kernel]      [k] free_pages_prepare
  0.00%    1  [kernel]      [k] free_pcpages_bulk
```

What is a precise event?

Precise Events

- Interrupt-based sampling is based on counting a specific performance event and waiting until it overflows. When an overflow happens, it takes a processor some time to stop the execution and tag the instruction that caused the overflow.
- The distance between the IP (instruction address) that caused the event to the IP where the event is tagged can be hundreds of processor instructions, causing confusion in performance analysis.
- With Intel PEBS (*Precise Event-Based Sampling*), the EventingIP field in the PEBS record indicates the instruction that caused the event.
- This is typically available only for a subset of supported events, called “Precise Events”.

Step 2: Locate the Place in the Code

Perf annotate locates the assembly line at which the bottleneck occurs

```
$ perf annotate --stdio -M intel foo
Percent | Disassembly of benchmark.exe for MEM_LOAD_RETIRE.L3_MISS
-----
: Disassembly of section .text:
:
: 0000000000400a00 <foo>:
: foo():

0.00 : 400a00: nop DWORD PTR [rax+rax*1+0x0]
0.00 : 400a08: nop DWORD PTR [rax+rax*1+0x0]
...
... # more NOPs
100.00 : 400e07: mov rax, QWORD PTR [rdi+rsi*1] <==
...
0.00 : 400e13: xor rax, rax
0.00 : 400e16: ret
```

Step 2: Locate the Place in the Code

The source code of the `main()` function:

- allocate an array bigger than the L3 cache.
- Generates a random index into array `a` and passes this index to the `foo()` function (in the `rdi` and `rsi` registers)

```
extern "C" { void foo(char* a, int n); }
const int _200MB = 1024*1024*200;
int main() {
    char* a = new char[_200MB]; // 200 MB buffer
    ...
    for (int i = 0; i < 100000000; i++) {
        int random_int = distribution(generator);
        foo(a, random_int);
    }
    ...
}
```

Step 2: Locate the Place in the Code

Perf annotate locates the assembly line at which the bottleneck occurs

```
$ perf annotate --stdio -M intel foo
Percent | Disassembly of benchmark.exe for MEM_LOAD_RETIRED.L3_MISS
-----
: Disassembly of section .text:
:
: 0000000000400a00 <foo>:
: foo():

0.00 : 400a00: nop DWORD PTR [rax+rax*1+0x0]
0.00 : 400a08: nop DWORD PTR [rax+rax*1+0x0]
...
# more NOPs
100.00 : 400e07: mov rax,QWORD PTR [rdi+rsi*1] <=
...
0.00 : 400e13: xor rax,rax
0.00 : 400e16: ret
```

a+random_int

Step 3: Fix the Issue

The `foo()` function executes some dummy work emulated with NOPs.

- This creates a time window between the computing of `random_int` and the moment when we load `* (a+random_int)`
- The presence of the time window allows us to start prefetching the memory location in parallel with the dummy work.

Step 3: Fix the Issue

The `foo()` function executes some dummy work emulated with NOPs.

- This creates a time window between the computing of `random_int` and the moment when we load `* (a+random_int)`
- The presence of the time window allows us to start prefetching the memory location in parallel with the dummy work.

```
for (int i = 0; i < 100000000; i++) {  
    int random_int = distribution(generator);  
    __builtin_prefetch(a+random_int, 0, 1);  
    foo(a, random_int);  
}
```

Step 3: Fix the Issue

The `foo()` function executes some dummy work emulated with NOPs.

- This creates a time window between the computing of `random_int` and the moment when we load `* (a+random_int)`
- The presence of the time window allows us to start prefetching the memory location in parallel with the dummy work.

```
for (int i = 0; i < 100000000; i++) {
    int random_int = distribution(generator);
+    __builtin_prefetch(a+random_int,0,1);
    foo(a, random_int);
}
```

- Memory prefetching decreases execution time from 8.5 seconds to 6.5 seconds.
- The number of `CYCLE_ACTIVITY.STALLS_L3_MISS` events becomes goes from 19 billion down to 2 billion.

TMA Remarks

Back-End Bound can be divided in

1. Memory bound:

- Performance limited by memory subsystem latency/bandwidth.

2. Core Bound:

- Performance limited by the CPU execution ports

TMA Remarks

Memory bound:

- The majority of un-tuned applications will be Memory bound.
- Most of the metrics identify which level of the memory hierarchy is the bottleneck. Optimizations should focus on moving data closer to the core.
- Store Bound can indicate dependencies - such as when loads in the pipeline depend on prior stores.

TMA Remarks

Core bound:

- Core Bound stalls are typically less common than memory bound.
- They occur when available computing resources are not sufficient. For example, a tight loop doing Floating Point (FP) arithmetic calculations on data that fits within cache. Some metrics helps to detect behaviors in this category:
 - the Divider metric identifies cycles when divider hardware is heavily used
 - the Port Utilization metric identifies competition for discrete execution units.

TMA Remarks

Front-End Bound can be reduced by:

- Changing code layout (co-locating hot code)
- Reducing branchy code
- code size optimization and compiler profile-guided optimization (PGO) are likely to reduce stalls in many cases.

Bad Speculation can be reduced through compiler techniques such as:

- Profile-Guided Optimization (PGO),
- avoiding indirect branches,
- eliminating error conditions that cause machine clears.

Correcting Bad Speculation issues may also help decrease the number of Front-End Bound stalls.

TMA Remarks

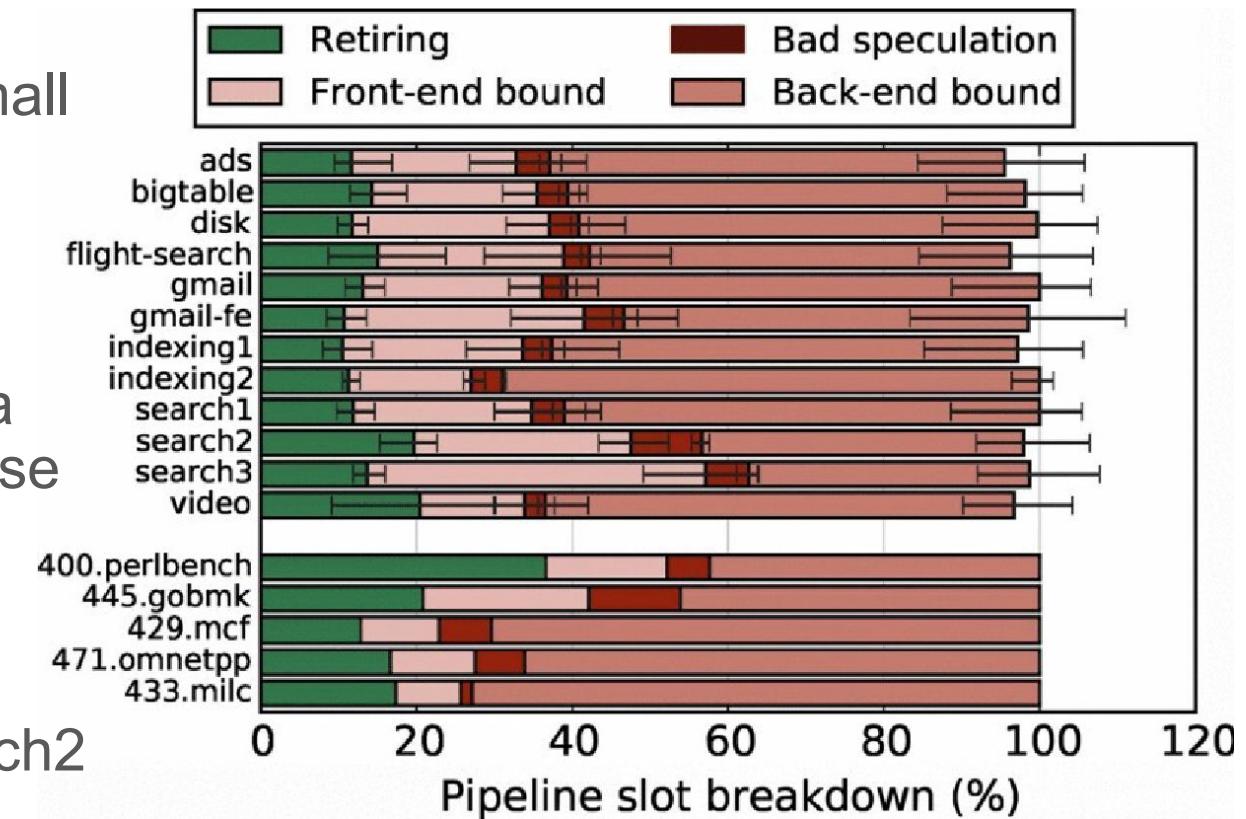
Ideally, after several iterations we would like to see the **Retiring** metric at 100%. Having 100% Retiring means a CPU is fully saturated and it crunches instructions at full speed.

- But it doesn't say anything about the quality of those instructions. A program can spin in a tight loop waiting for a lock; that would show a high Retiring metric, but it doesn't do any useful work.
- If you see a high Retiring metric but slow overall performance probably the program has a hotspot that was not vectorized.
- Another performance issue that will fall under the retiring category is heavy use of the micro-sequencer (e.g. to manage subnormal FP)

TMA Remarks

Achieving 100% Retiring in real-world applications is rare.

- Most data center workloads have a very small fraction in the Retiring bucket.
- BackendBound is the primary source of performance issues.
- FrontendBound is a bigger problem for data center workloads than in SPEC2006 because those applications typically have large codebases with poor locality.
- Finally, some workloads suffer from branch mispredictions more than others, e.g., search2 and 445.gobmk.



TMA breakdown of Google's datacenter workloads along with several SPEC CPU2006 benchmarks,

TMA with Intel Vtune profiler

Linux `perf` is a lightweight, command-line profiling tool:

- available on all Linux systems
- Low overhead,
- easy to integrate in scripts
- limited visualization (requires external tools)
- No deep microarchitecture insight (limited to events counters)

Intel Vtune:

- More detailed microarchitecture-level metrics
- Detailed visualization
- Provide optimization hints

TMA with Intel Vtune profiler

Microarchitecture Exploration Microarchitecture Exploration ▾ ⓘ ⓘ

Analysis Configuration Collection Log Summary Bottom-up Event Count Platform

⌚ Elapsed Time: 38.454s >

Clockticks:	171,490,200,000
Instructions Retired:	204,835,200,000
CPI Rate:	0.837
⌚ Retiring:	43.5% of Pipeline Slots
⌚ Light Operations:	28.0% of Pipeline Slots
⌚ Heavy Operations:	15.5% of Pipeline Slots
⌚ Few Uops Instructions:	15.0% of Pipeline Slots
⌚ Microcode Sequencer:	0.5% of Pipeline Slots
⌚ Front-End Bound:	0.5% of Pipeline Slots
⌚ Bad Speculation:	0.5% of Pipeline Slots
⌚ Back-End Bound:	55.5% of Pipeline Slots
⌚ Memory Bound:	21.3% of Pipeline Slots
⌚ L1 Bound:	1.5% of Clockticks
⌚ L2 Bound:	3.1% of Clockticks
⌚ L3 Bound:	14.2% of Clockticks
⌚ Contested Accesses:	0.0% of Clockticks
⌚ Data Sharing:	0.0% of Clockticks
⌚ L3 Latency:	100.0% of Clockticks
⌚ SQ Full:	0.9% of Clockticks
⌚ DRAM Bound:	3.4% of Clockticks
⌚ Store Bound:	0.0% of Clockticks
⌚ Core Bound:	34.2% of Pipeline Slots
⌚ Divider:	0.0% of Clockticks
⌚ Port Utilization:	35.2% of Clockticks
Average CPU Frequency:	4.7 GHz
Total Thread Count:	2
Paused Time:	1.376s

⌚ Effective Physical Core Utilization: 12.2% (0.980 out of 8) >

Effective Logical Core Utilization: 6.2% (0.986 out of 16) ⓘ

The μPipe diagram illustrates inefficiencies in CPU usage. It consists of four colored segments representing different performance metrics: Memory Bound (orange), Retiring (green), Core Bound (orange), and a final green segment. The segments are stacked vertically, with the width of each segment representing its contribution to the total CPU usage. The segments are separated by black lines, creating a stepped, pipe-like appearance. The segments are labeled with their respective percentages and descriptions.

21.3% - Memory Bound
The metric value is high. This can indicate that the significant fraction of

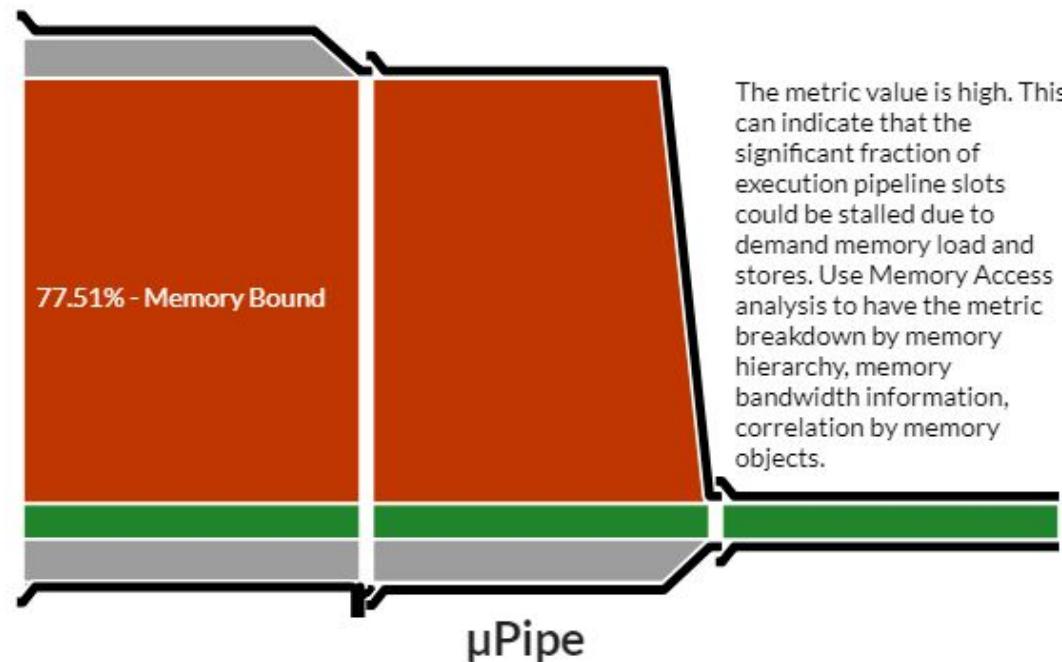
43.5% - Retiring

34.2% - Core Bound
This metric represents how much Core non-memory issues were of a bottleneck. Shortage in hardware compute resources or

This diagram represents inefficiencies in CPU usage. Treat it as a pipe with an output flow equal to the "pipe efficiency" ratio: (Actual Instructions Retired)/(Maximum Possible [Instruction Retired](#)). If there are pipeline stalls decreasing the pipe efficiency, the pipe shape gets more narrow.

Microarchitecture Pipe

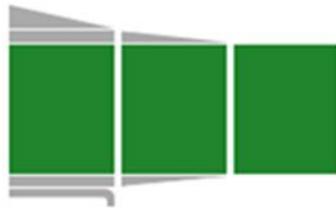
The Microarchitecture Pipe graphically identifies inefficiencies in the CPU utilization.



The output flow is the ratio: **Actual Instructions Retired/Possible Maximum Instruction Retired** (pipe efficiency). If there are pipeline stalls decreasing retiring, the pipe shape gets narrow.

Microarchitecture Pipe examples

A. 100% retiring



B. Memory bound execution



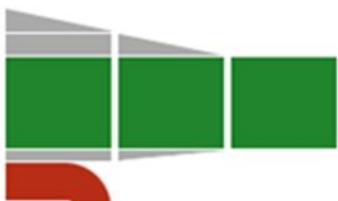
C. Core bound execution



D. Front End bound execution



E. Bad speculation

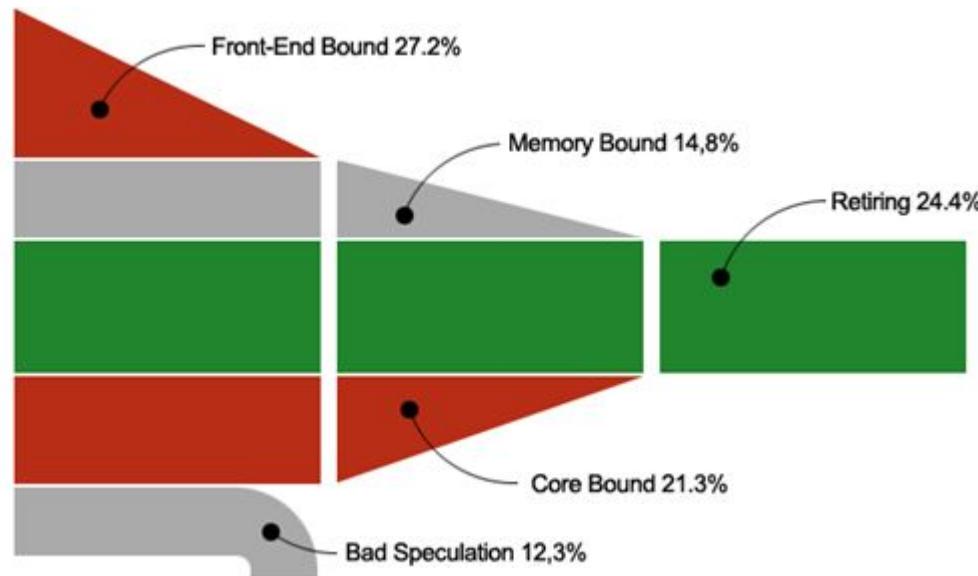


F. Memory and Bad speculation



Microarchitecture Pipe Example

This is an example of a pipe representing significant **Front-End Bound** and **Core Bound** issues limiting the whole efficiency to 24.4%:



Case Study: n-body problem

We already optimized the n-body using vectorization

```
// Compute pairwise forces
for (int i = 0; i < Np; i++) {
    for (int j = i+1; j < Np; j++) {
        double dx = p[j].x - p[i].x;
        double dy = p[j].y - p[i].y;
        double dz = p[j].z - p[i].z;
        [...]
        double inv_r3 = 1.0/r2 *sqrt(r2);
    }
}
```

Now we check perform TMA on the nbody program

TMA on the nbody code

Intel VTune Profiler (maestrale)

Project Navigator + Welcome x r013ue x

Microarchitecture Exploration Microarchitecture Exploration Bottom-up Event Count Platform

Elapsed Time : 1.829s

Clockticks: 5,383,150,000
Instructions Retired: 10,837,600,000
CPI Rate: 0.497

Bottlenecks View (preview):

Mispredictions	0.0%	of Pipeline Slots
Big Code	0.0%	of Pipeline Slots
Instruction Fetch Bandwidth	N/A*	of Pipeline Slots
Cache Memory Bandwidth	N/A*	of Pipeline Slots
Compute Bound Estimation	46.4% ↘	of Pipeline Slots
Irregular Overhead	N/A*	of Pipeline Slots
Branching Overhead	1.1%	of Pipeline Slots
Useful Work	27.0% ↘	of Pipeline Slots

Top-down Microarchitecture Analysis (TMA):

Retiring	28.1%	of Pipeline Slots
Front-End Bound	3.6%	of Pipeline Slots
Bad Speculation	6.7%	of Pipeline Slots
Back-End Bound	61.5% ↘	of Pipeline Slots
Memory Bound	5.0%	of Pipeline Slots
Core Bound	56.5% ↘	of Pipeline Slots
Divider	89.5% ↘	of Clockticks
FP Divider	89.1% ↘	of Clockticks
INT Divider	0.5%	of Clockticks
Serializing Operations	0.7%	of Clockticks
AMX Busy	0.0%	of Clockticks
Port Utilization	38.3% ↘	of Clockticks
Cycles of 0 Ports Utilized	0.0%	of Clockticks
Cycles of 1 Port Utilized	24.8% ↘	of Clockticks
Cycles of 2 Ports Utilized	27.6% ↘	of Clockticks
Cycles of 3+ Ports Utilized	38.1%	of Clockticks
Vector Capacity Usage (FPU)	16.2% ↘	

Average CPU Frequency: 3.0 GHz
Total Thread Count: 2
Paused Time: 0s

*N/A is applied to metrics with undefined value. There is no data to calculate the metric.

Effective CPU Utilization: 6.1% ↘

Average Effective CPU Utilization: 0.977 out of 16

μPipe

This diagram represents inefficiencies in CPU usage. Treat it as a pipe with an output flow equal to the "pipe efficiency" ratio: (Actual Instructions Retired)/(Maximum Possible Instruction Retired). If there are pipeline stalls decreasing the pipe efficiency, the pipe shape gets more narrow.

TMA on the nbody code

Intel VTune Profiler (maestrale)

Project Navigator + Welcome r013ue

Microarchitecture Exploration Microarchitecture Exploration Analysis Configuration Collection Log Summary Bottom-up Event Count Platform

Elapsed Time : 1.829s

Clockticks: 5,383,150,000
Instructions Retired: 10,837,600,000
CPI Rate: 0.497

Bottlenecks View (preview):

Mispredictions:	0.0% of Pipeline Slots
Big Code:	0.0% of Pipeline Slots
Instruction Fetch Bandwidth:	N/A* of Pipeline Slots
Cache Memory Bandwidth:	N/A* of Pipeline Slots
Compute Bound Estimation:	46.4% of Pipeline Slots
Irregular Overhead:	N/A* of Pipeline Slots
Branching Overhead:	1.1% of Pipeline Slots
Useful Work:	27.0% of Pipeline Slots

Top-down Microarchitecture Analysis (TMA):

Retiring:	28.1% of Pipeline Slots
Front-End Bound:	3.6% of Pipeline Slots
Bad Speculation:	6.7% of Pipeline Slots
Back-End Bound:	61.5% of Pipeline Slots
Memory Bound:	5.0% of Pipeline Slots
Core Bound:	56.5% of Pipeline Slots
Divider:	89.5% of Clockticks
FP Divider:	89.1% of Clockticks
INT Divider:	0.5% of Clockticks
Serializing Operations:	0.7% of Clockticks
AMX Busy:	0.0% of Clockticks
Port Utilization:	38.3% of Clockticks
Cycles of 0 Ports Utilized:	0.0% of Clockticks
Cycles of 1 Port Utilized:	24.8% of Clockticks
Cycles of 2 Ports Utilized:	27.6% of Clockticks
Cycles of 3+ Ports Utilized:	38.1% of Clockticks
Vector Capacity Usage (FPU):	16.2% of Clockticks

Average CPU Frequency: 3.0 GHz
Total Thread Count: 2
Paused Time: 0s

*N/A is applied to metrics with undefined value. There is no data to calculate the metric.

Effective CPU Utilization: 6.1%

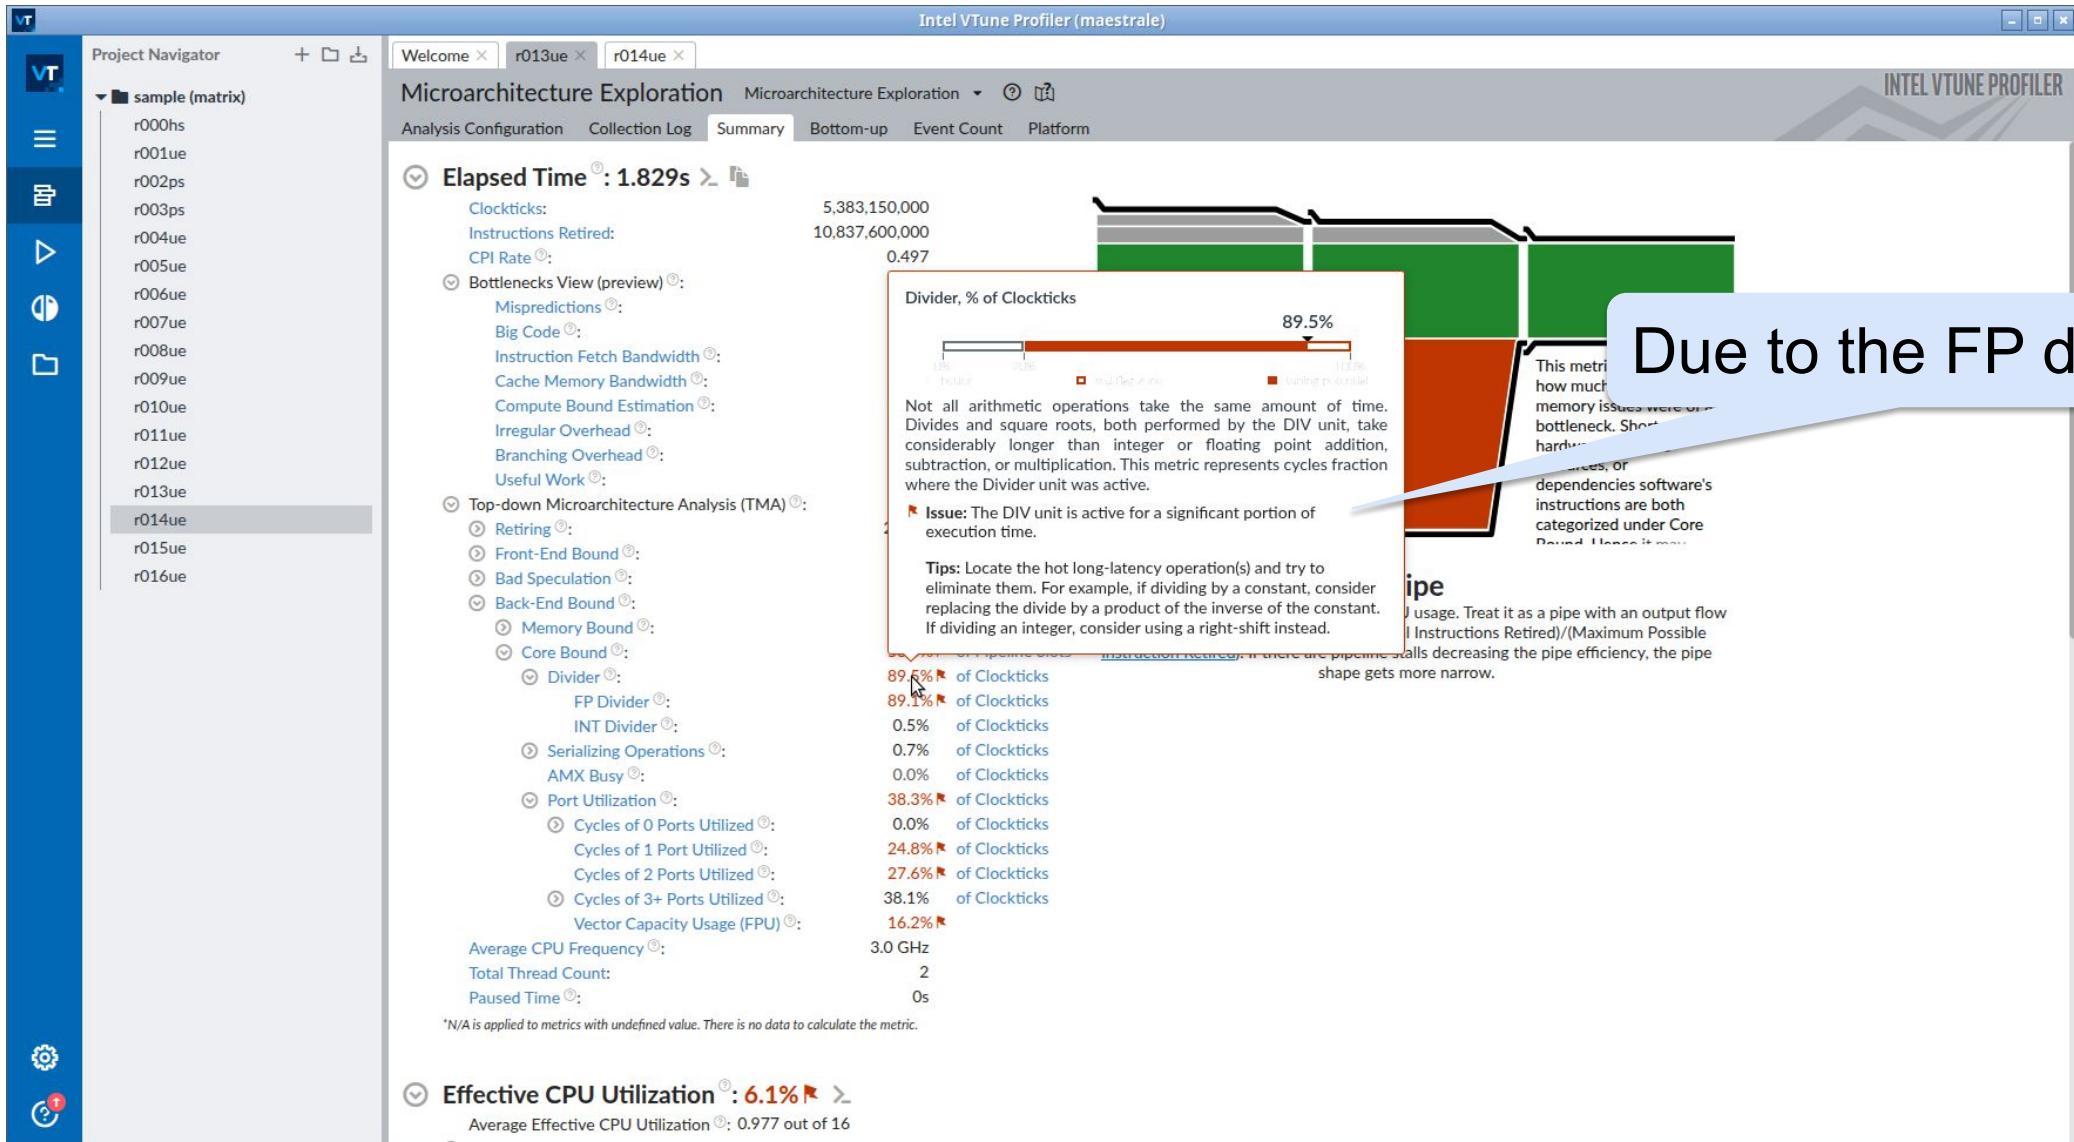
Average Effective CPU Utilization: 0.977 out of 16

μPipe

This diagram represents inefficiencies in CPU usage. Treat it as a pipe with an output flow equal to the "pipe efficiency" ratio: (Actual Instructions Retired)/(Maximum Possible Instruction Retired). If there are pipeline stalls decreasing the pipe efficiency, the pipe shape gets more narrow.

It is core bounded

TMA on the nbody code



RSQRT — Reciprocal Square Root Approximation

- The `rsqrt` computes an **approximate value of $1 / \sqrt{x}$** . It has 14-bit precision

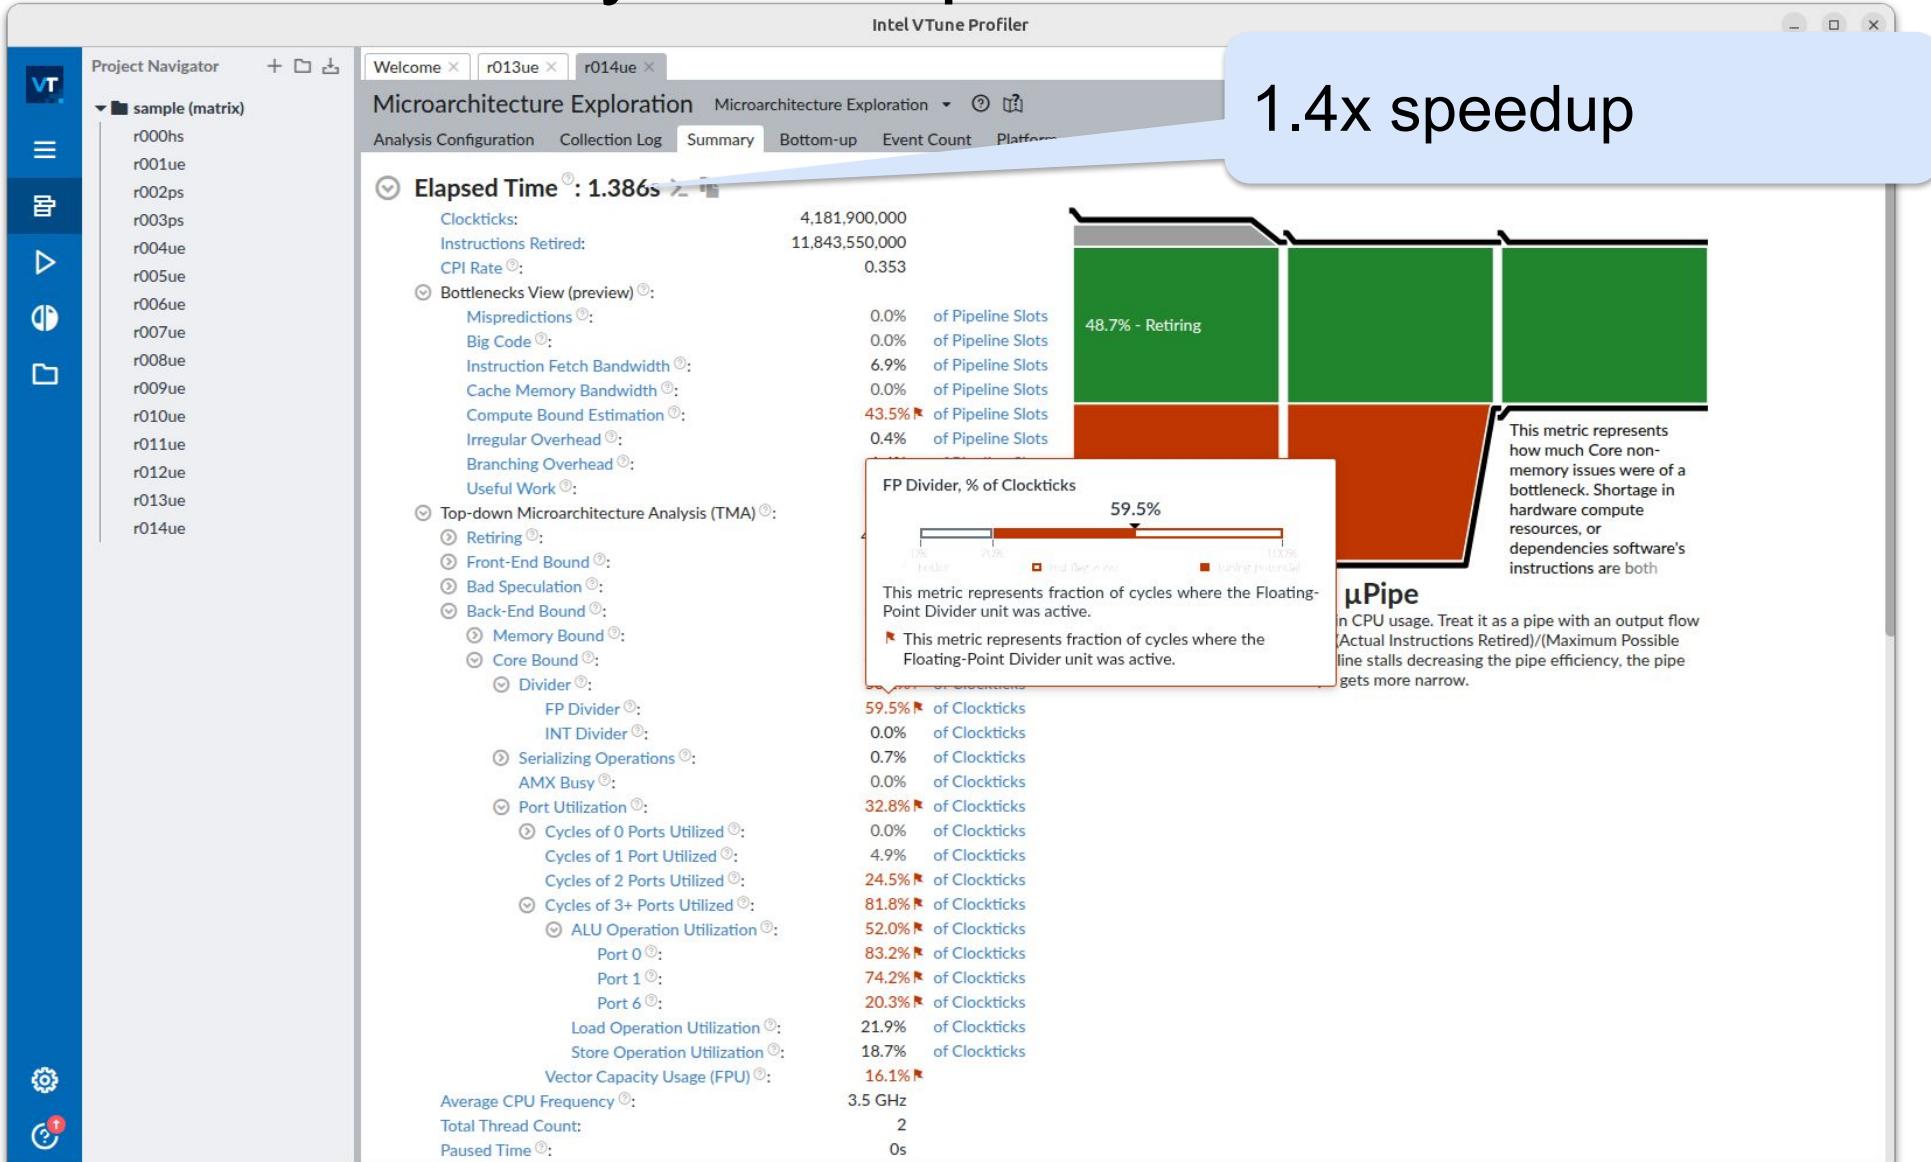
```
RSQRTSS xmm1, xmm2 ; xmm1[i] = 1 / sqrt(xmm2[i]) (approx)
```

- It exists a vectorized AVX-512 instruction

```
vrsqrt14ps zmm1, zmm2 ; zmm1[i] = 1 / sqrt(zmm2[i]) (approx)
```

- Typically followed by a Newton–Raphson refinement for full precision if needed.

TMA on the nbody with rsqrt

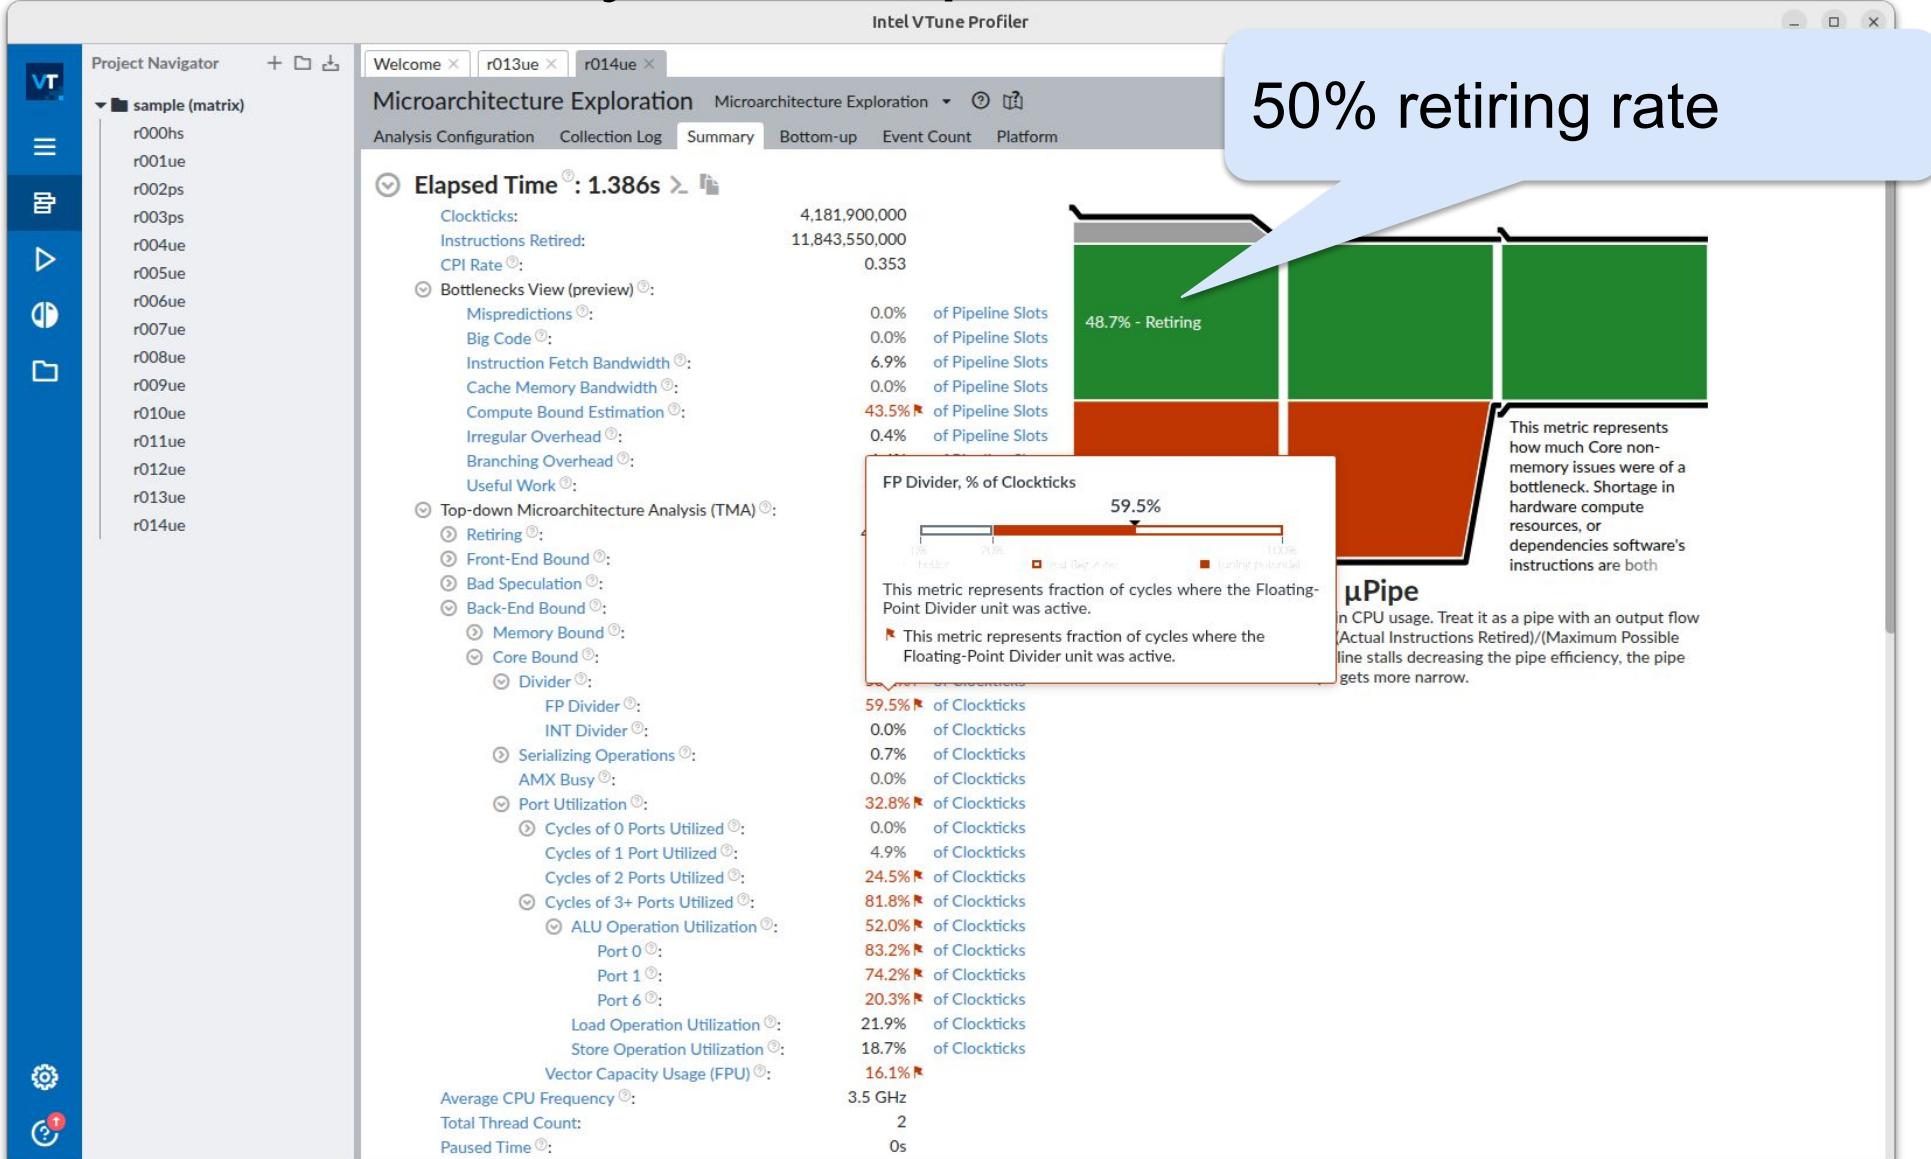


TMA on the nbody with rsqrt



Less usage of divider

TMA on the nbody with rsqrt



Bibliography

*Denis, Bakhvalov, Performance Analysis And Tuning On Modern CPUs:
Second Edition*

Intel VTune Profiler Performance Analysis Cookbook