

# Optimizing Computations

# Optimizing Computations

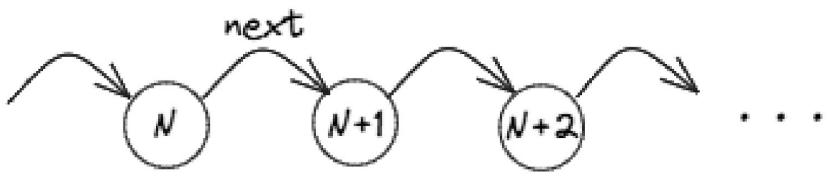
When the TMA methodology is applied, inefficient computations are usually reflected in the **Core Bound** and, to some extent, in the **Retiring** categories.

There are two main issues for **Core Bound** categories:

- 1) **Data dependencies.** For example, a long sequence of dependent operations may lead to low ILP and wasting many execution slots.
- 2) Shortage in hardware computing resources. This indicates that certain execution units are overloaded (also known as *execution port contention*). For example, scientific applications may run many divisions and square root operations. However, there is a limited number of multipliers and dividers in any given CPU core.
  - When port contention occurs, instructions queue up waiting for their turn to be executed. This type of performance bottleneck is very specific to a particular CPU microarchitecture and usually doesn't have a cure.

# Data Dependencies

- When a program statement refers to the output of a preceding statement, we say that there is a *data dependency* between the two statements.
- We also use the terms *dependency chain* or *data flow dependencies*.



```
while (n) {  
    sum += n->val;  
    n = n->next;  
}
```

- Traversing a linked list is a typical example

# Data Dependencies

- Long data dependencies forces CPU to execute code sequentially, defeating the main advantage of modern superscalar CPUs.
  - For example, pointer chasing doesn't benefit from OOO execution and thus will run at the speed of an in-order CPU.
- Dependency chains are a major source of performance bottlenecks.
- We want to find a critical dependency chain in a hot piece of code, such as a loop or function.

# Data Dependency example

Example with Random Particle Motion. *Where is the dependency chain?*

```
struct Particle {
    float x; float y; float
velocity;
};

class XorShift32 {
    uint32_t val;
public:
    XorShift32 (uint32_t seed) :
val(seed) {}
    uint32_t gen() {
        val ^= (val << 13);
        val ^= (val >> 17);
        val ^= (val << 5);
    return val;
}
}
```

```
/* Map degrees [0;UINT32_MAX) to radians
[0;2*pi)*/
float DEGREE_TO_RADIAN = (2 * PI_D) / UINT32_MAX;
void particleMotion(vector<Particle>
&particles,uint32_t seed) {
    XorShift32 rng(seed);
    for (int i = 0; i < STEPS; i++)
        for (auto &p : particles) {
            uint32_t angle = rng.gen();
            float angle_rad = angle * DEGREE_TO_RADIAN;
            p.x += cosine(angle_rad) * p.velocity;
            p.y += sine(angle_rad) * p.velocity;
        }
}
```

# Dependency chains

**Rule of thumb: if an instruction is on a critical path, look at its latency, otherwise look at its throughput.**

On every loop iteration, we have 5 fmul and 4 fmadd instructions. (outside the critical path)

The gen() function requires 6 cc

Since the M1 processor can run 4 instructions per cycle we will have at least  $9/4 = 2.25$  cycles to issue all the fp instructions. So, we have two performance limits:

- 1) 6 cycles per iteration due to the dependency chain (SW imposed)
- 2) 2.25 cycles per iteration (HW imposed)

# Break dependency chain

- We can employ an additional RNG object to break the chain
- Running “before” and “after” versions
  - running time go from 19ms to 10ms
  - IPC goes from 4.0 to 7.1

```
void particleMotion(vector<Particle>& particles, uint32_t seed) {  
    XorShift32 rng1(seed1);  
    XorShift32 rng2(seed2);  
    for (int i = 0; i < STEPS; i++)  
        for (int j=0; j< particles.size(); j +=2) {  
            uint32_t angle1 = rng1.gen();  
            uint32_t angle2 = rng2.gen();  
            particles[j] ...  
            particles[j+1] ...  
        }  
}
```

# Inlining Functions

A function call require to pass arguments, preserve registers and manage return values.

- Usually functions uses a series of PUSH instructions (a prologue) when the function start, and a series of POP instructions (an epilogue) when the function return.
- When a function is small, the overhead of calling a function (prologue and epilogue) can be very pronounced. This overhead can be eliminated by inlining a function body into the place where it is called.
- The primary mechanism for function inlining in many compilers relies on a cost model. Inlining happens if the cost is less than a threshold, which is usually a fixed number;

# Inlining Functions

Example of functions candidate for inlining

- Tiny functions (wrappers) are almost always inlined.
- Functions with a single call site are preferred candidates for inlining.

Instead, large functions usually are not inlined as they bloat the code of the caller function.

# Inlining Functions

Profiling data can help in selecting functions to inline

Overhead (%)	Source code & Disassembly
3.77	: 418be0: push r15 <b># prologue</b>
4.62	: 418be2: mov r15d,0x64
2.14	: 418be8: push r14
1.34	: 418bea: mov r14,rsi
3.43	: 418bed: push r13
3.08	: 418bef: mov r13,rdi
1.24	: 418bf2: push r12
1.14	: 418bf4: mov r12,rcx
3.08	: 418bf7: push rbp
3.43	: 418bf8: mov rbp,rdx
1.94	: 418bfb: push rbx
0.50	: 418bfc: sub rsp,0x8
...	
	# function body
...	
4.17	: 418d43: add rsp,0x8 <b># epilogue</b>
3.67	: 418d47: pop rbx
0.35	: 418d48: pop rbp
0.94	: 418d49: pop r12
4.72	: 418d4b: pop r13
4.12	: 418d4d: pop r14
0.00	: 418d4f: pop r15
1.59	: 418d51: ret

# Loop Optimizations

Loops are at the heart of nearly all high-performance programs. Small changes in such a critical piece of code may have a large impact on the performance of a program.

Several techniques to improve loop performance exist:

- *Low-level Optimizations*: usually apply to a single loop and can be done by compiler.
- *High-level Optimizations*: usually apply to nested loop and require must be done manually

# Low-level Optimizations

Low level optimizations transform the code inside a single loop. Generally, compilers are good at doing such transformations.

- *Loop Invariant Code Motion (LICM)*: when a value doesn't change across loop iterations, we can move a loop invariant expression outside of the loop.
- *Loop Unrolling*: we can unroll a loop and perform multiple iterations for each increment of the induction variable
- *Loop Strength Reduction (LSR)*: replace expensive instructions with cheaper ones. Such transformation can be applied to all expressions that use an induction variable (see the sieve exercise).

```
for (int i = 0; i < N; ++i)
    a[i] = b[i * 10] * c[i];
```



```
for (int i = 0; i < N; ++i) {
    a[i] = b[j] * c[i];
    j += 10;
}
```

# High-level Optimizations

High-level loop transformations change the structure of loops and often affect multiple nested loops.

- From a compiler perspective, it is very difficult to prove the legality of such transformations and justify their performance benefit.
- Examples of optimization are:
  - Loop Interchange
  - Loop Blocking (Tiling)
  - Loop Fusion and Distribution (Fission):
  - Loop Unroll and Jam

# Loop Interchange

Loop Interchange exchanges the loop order of nested loops

```
for (int i = 0; i < N; i++)
    for (int j = 0; j < N; j++)
        a[i][j] += b[i][j]*c[i][j];
```



```
for (int j = 0; i < N; i++)
    for (int i = 0; j < N; j++)
        a[i][j] += b[i][j]*c[i][j];
```

- Loop interchange performs sequential memory accesses to the elements of a multi-dimensional array, improving the spatial locality of memory accesses. The transformation helps to eliminate memory bandwidth and memory latency bottlenecks.
- Loop Interchange require that loops are *perfectly nested*. Interchanging imperfect loop nests require code restructuring.

# Converting imperfectly-nested loops

Example 1: Isolating statements that prevent perfect loop nesting

```
for (int i = 0; i < N; i++) {  
    a[i]=0.0;  
    for (int j = 0; j < N; j++)  
        a[i] += b[i][j]*c[i][j];  
}
```



```
for (int i = 0; i < N; i++)  
    a[i]=0.0;  
for (int j = 0; j < N; j++)  
    for (int i = 0; i < N; i++)  
        a[i] += b[i][j]*c[i][j];
```

# Converting imperfectly-nested loops

Example 2: promoting temporary scalars to vectors

```
for (int i = 0; i < N; i++) {  
    double sum = 0.0;  
    double dot = 0.0;  
    for (int j = 0; j < N; j++) {  
        sum += b[i][j]+c[i][j];  
        dot += b[i][j]*c[i][j];  
    }  
    a[i]= dot/sum;  
}
```

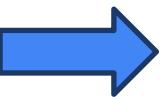


```
double sum_array [N];  
double dot_array [N];  
for (int i = 0; i < N; i++) {  
    sum_array [i]=0.0;  
    dot_array [i]=0.0;  
}  
for (int i = 0; i < N; i++) {  
    for (int j = 0; j < N; j++) {  
        sum_array[i] += b[i][j]+c[i][j];  
        dot_array[i] += b[i][j]*c[i][j];  
    }  
}  
for (int i= 0; i < N; i++)  
    a[i]= dot_array[i]/sum_array[i];  
}
```

# Loop Blocking (Tiling)

The idea of tiling is to split the multidimensional execution range into smaller chunks (blocks or tiles) so that each block will fit in the CPU caches.

```
// linear traversal
for (int i = 0; i < N; i++)
    for (int j = 0; j < N; j++)
        a[i][j] += b[j][i];
```



```
// traverse in 8*8 blocks
for (int ii = 0; ii < N; ii+=8)
    for (int jj = 0; jj < N; jj+=8)
        for (int i = ii; i < ii+8; i++)
            for (int j = jj; j < jj+8; j++)
                a[i][j] += b[j][i];
// remainder (not shown)
```

- Typically, engineers optimize a tiled algorithm for the size of caches that are private to each CPU core. Since private caches changes from generation to generation hardcoding a block size presents its own set of challenges.

# Loop Fusion and Distribution (Fission)

Separate loops can be fused when they iterate over the same range and **do not reference each other's data**.

```
for (int i = 0; i < N; i++)
    sum += b[i]+c[i];
for (int i = 0; i < N; i++)
    dot += b[i]*c[i];
```



```
for (int i = 0; i < N; i++) {
    sum += b[i]+c[i];
    dot += b[i]*c[i];
}
```

- **Loop Fusion** helps to reduce the loop overhead and to improve the temporal locality of memory accesses.

# Loop Fusion and Distribution (Fission)

The opposite choice is called loop distribution/fission

```
for (int i = 0; i < N; i++) {  
    sum += a[i]+b[i];  
    dot += c[i]*d[i];  
}
```



```
for (int i = 0; i < N; i++)  
    sum += a[i]+b[i];  
for (int i = 0; i < N; i++)  
    dot += c[i]*d[i];
```

- **Loop Fission** helps in situations with a high cache contention and reduces register pressure.
- It also improves instruction cache utilization.
- Finally, when distributed, each small loop can be further optimized separately by the compiler.

# Loop Unrolling

**Loop unrolling** increases the Instruction-Level Parallelism of the inner loop since more independent instructions are executed inside the inner loop.

```
for (int i = 0; i < N; i++) {  
    diff += a[i]-b[i];  
}
```



```
for (int i = 0; i < N; i+=2) {  
    diff1 += a[i]-b[i];  
    diff2 += a[i+1]-b[i+1];  
}  
diff = diff1 + diff2;
```

- Unrolling breaks the dependency chain `diff` in the initial variant

# Discovering Loop Optimization Opportunities

- Loop optimization opportunities can be identified looking into the **Compiler optimization reports** or examining the machine code.
- Many compiler **pragmas** are available to control loop transformations (loop vectorization, loop unrolling, loop distribution, etc. ).
- TMA and roofline model can help to assessment loop performance. Performance is limited by one or many of the following factors:
  - memory latency,
  - memory bandwidth,
  - computing capabilities
- Once the performance bottlenecks in a loop have been identified, try applying the required transformations previously discussed.