


# High-Performance Computing

Daniele De Sensi

# Software Stack



# DISCLAIMER

Most of the following content is CUDA-related, but the same would apply to AMD GPUs. Also, regardless of the specific technology, those techniques (overlap etc...) are general and would apply to different accelerators etc...

# Multi GPU Programming in MPI

Slides adapted from slides from:  
Lena Oden, Fern Universität in Hagen  
Simon Garcia de Gonzalo, Sandia National Laboratories  
<https://github.com/FZJ-JSC/tutorial-multi-gpu/>

# Multi-GPU Computing: What you will learn

- CUDA-aware MPI
- Example: Jacobi-Solver
- Under the hood (why you should use CUDA-aware MPI)
  - GPUs in Clusters
  - CUDA Unified Virtual Addressing
  - GPUDirect P2P and GPUDirect RDMA

# Message Passing Interface - MPI

- Standard to exchange data between processes via messages
  - Defines API to exchange messages
    - Point to Point: e.g. MPI\_Send, MPI\_Recv
    - Collectives: e.g. MPI\_Reduce, MPI\_Allreduce, MPI\_Bcast
- Multiple implementations (open source and commercial)
  - Bindings for C/C++, Fortran, Python, ...
  - e.g. MPICH, OpenMPI, MVAPICH, IBM Spectrum MPI, Cray MPT, ParaStation MPI, ...

# Example MPI Program

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    // Initialize the MPI environment
    MPI_Init(NULL, NULL);
    int size;
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
    MPI_Finalize();
}
```

```
mpicc -o hello_mpi.out hello_mpi.c
mpirun -n 4 ./hello_mpi.out
```

# CUDA-aware MPI

CUDA-aware MPI allows you to use Pointers to GPU-Memory as source and destination

```
//MPI rank 0  
MPI_Send(s_buf_d, n, MPI_BYTE, size-1, tag, MPI_COMM_WORLD);  
  
//MPI size-1  
MPI_Recv(r_buf_d, n, MPI_BYTE, 0, tag, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
```

Pointer to  
GPU Memory

# LAUNCH MPI+CUDA


- Launch one process per GPU
- How to use CUDA-aware MPI
  - *MVAPICH*: **\$MV2\_USE\_CUDA=1** mpirun -np \${np} ./myapp <args>
  - *Open MPI*: *CUDA-aware features are enabled per default (using UCX)*
  - *Cray*: MPICH\_RDMA\_ENABLED\_CUDA
  - *IBM Spectrum MPI*: \$mpirun -gpu -np \${np} ./myapp <args>
  - *ParaStation MPI* (using UCX): **\$PSP\_CUDA=1** mpirun -np \${np} ./myapp <args>

# How to compile

```
nvcc -o my_kernel.o $(NVCC_FLAGS) my_kernels.cu -c  
mpicc -o my_multiGPUapp -lcudart my_kernel.o my_multiGPUapp.c
```

# Example: Jacobi Solver


- Solves the 2D-Laplace Equation on a rectangle
- $\Delta \mathbf{u}(x, y) = \mathbf{0} \forall (x, y) \in \Omega \setminus \delta\Omega$
- Dirichlet boundary conditions on left/right boundaries (constant values on boundaries)
- Reflecting boundaries on top and bottom
- Iterative solver:  $u(t+1) = f(u(t))$



# Example: Jacobi solver

While not converged do Jacobi step:

```
int iy = blockIdx.y * blockDim.y + threadIdx.y
int ix = blockIdx.x * blockDim.x + threadIdx.x
if (iy < ny-1 && ix < nx-1) {
    new_val =
        0.25 * (a[iy * nx + ix + 1] + a[iy * nx + ix - 1
                                              a[(iy + 1) * nx + ix] + a[(iy - 1
                                                * nx + ix]);
    a_new[iy * nx + ix] = new_val;
}
```





apply boundary Condition


swap a\_new and a

next iteration

# Domain Decomposition



- \* Minimize number of neighbors: Communicate to fewer neighbors
- \* Optimal for latency bound communication,
- \* Contiguous Transfers




- \* Minimize surface area/volume ratio: Communicate less data
- \* Optimal for bandwidth bound communication
- \* Non-Contiguous Transfers<sup>10</sup>

# Jacobi example: Top and Bottom Boundaries



```
MPI_Sendrecv(a_new_d+offset_last_row, m-2, MPI_DOUBLE,
             b_nb, 1, a_new_d+offset_top_boundary, m-2,
             MPI_DOUBLE, t_nb-1, MPI_COMM_WORLD,
             MPI_STATUS_IGNORE);

MPI_Sendrecv(a_new_d+offset_first_row, m-2, MPI_DOUBLE,
             t_nb, 0, a_new_d+offset_bottom_boundary, m-2,
             MPI_DOUBLE, b_nb, 0, MPI_COMM_WORLD,
             MPI_STATUS_IGNORE);
```




# Jacobi: Left (and right) neighbours


```
//right neighbor omitted  
pack<<<gs,bs,0,s>>>(to_left_d, u_new_d, n, m);  
cudaStreamSynchronize(s);  
MPI_Sendrecv(to_left_d, n-2, MPI_DOUBLE, l_nb, 0,  
             from_right_d, n-2, MPI_DOUBLE, r_nb, 0,  
             MPI_COMM_WORLD, MPI_STATUS_IGNORE );  
unpack<<<gs,bs,0,s>>>(u_new_d, from_right_d, n, m);  
cudaStreamSynchronize(s);
```



# Process Mapping on Multi GPU Systems: One GPU per Process



# Process Mapping on Multi GPU Systems: One GPU per Process




# Distribute GPUs to local Nodes


```
MPI_Comm local_comm;  
MPI_Comm_split_type(MPI_COMM_WORLD, MPI_COMM_TYPE_SHARED, rank,  
                    MPI_INFO_NULL, &local_comm);  
int local_rank = -1;  
MPI_Comm_rank(local_comm,&local_rank);  
MPI_Comm_free(&local_comm);  
  
int num_devs = 0;  
cudaGetDeviceCount(&num_devs);  
cudaSetDevice(local_rank%num_devs);
```

# CUDA Unified Virtual Addressing

- One address space for all CPU and GPU memory
  - Determine physical memory location from a pointer value
  - Enable libraries to simplify their interfaces (e.g. MPI )
- Supported on devices with compute capability 2.0+ for
  - 64-bit applications on Linux and Windows (+TCC)




# Basics: GPUDirect




HCA = Host Channel Adapter = NIC (Network Card Interface)

# CUDA-aware MPI with GPUDirect RDMA




```
MPI_Send(s_buf_d, size, MPI_BYTE, 1, tag, MPI_COMM_WORLD);
MPI_Recv(r_buf_d, size, MPI_BYTE, 0, tag, MPI_COMM_WORLD, &stat);
```

# CUDA-aware MPI without GPUDirect RDMA



```
MPI_Send(s_buf_d, size, MPI_BYTE, 1, tag, MPI_COMM_WORLD);
MPI_Recv(r_buf_d, size, MPI_BYTE, 0, tag, MPI_COMM_WORLD, &stat);
```

# Regular MPI




```
cudaMemcpy(s_buf_h,s_buf_d,size,cudaMemcpyDeviceToHost);
MPI_Send(s_buf_h,size,MPI_BYTE,1,tag,MPI_COMM_WORLD);

MPI_Recv(r_buf_h,size,MPI_BYTE,0,tag,MPI_COMM_WORLD,&stat);
cudaMemcpy(r_buf_d,r_buf_h,size,cudaMemcpyHostToDevice);
```


# Performance Results GPUDirect RDMA

OpenMPI 4.1.0RC1 + UCX 1.9.0 on JUWELS Booster



# Performance Results GPUDirect P2P

OpenMPI 4.1.0RC1 + UCX 1.9.0 on JUWELS-Booster



# Summary

- CUDA-aware MPI allows efficient communication for multi-GPU applications
  - Allows MPI-communication operations from GPU memory buffers
  - Simplified programming
  - Use GPUDirect technologies for performance
  - Minimizes data copies
- Most MPI versions have support for CUDA-aware MPI
- Best practice: One process per rank
  - Use local communicator in MPI

# Hands-On

# Instructions

- Join the assignment on Github Classroom: <https://classroom.github.com/a/D9grIARs>
  - At the moment, there is no auto-grading yet, check correctness by yourself
- Access the cluster and clone the repository
- Fill the «TODOs» in the jacobi.cu code
- Run “source /home/guest/init-hpc.sh” to load the MPI, CUDA compilers etc...
- Run «make run» to compile and run the code (by default on 4 GPUs, one GPU per node)

# Solution

```
int num_devices = 0;
// TODO: Get the available GPU devices into `num_devices` and use it and the current rank to set the active GPU.
CUDA_RT_CALL(cudaGetDeviceCount(&num_devices));
CUDA_RT_CALL(cudaSetDevice(local_rank%num_devices));
CUDA_RT_CALL(cudaFree(0));

//TODO: Compute top and bottom neighbor, use reflecting/periodic boundaries. This means rank 0 and rank (size-1)
// exchange data
const int top = rank > 0 ? rank - 1 : (size - 1);
const int bottom = (rank + 1) % size;
```

# Solution

```
// TODO: Use MPI_Sendrecv to exchange the data with the top and bottom neighbors
// Use CUDA-aware MPI here, i.e. receive the data directly in a_new on the GPU and send it from there
// without manually copying the data.

// The first newly calculated row ('iy_start') is sent to the top neighbor and the bottom boundary row
// (`iy_end`) is received from the bottom process.
// The last calculated row (`iy_end-1`) is sent to the bottom process and the top boundary (`0`) is received from the top
// Don't forget to synchronize the computation on the GPU before starting the data transfer
CUDA_RT_CALL(cudaEventSynchronize(compute_done));
PUSH_RANGE("MPI", 5)
MPI_CALL(MPI_Sendrecv(a_new + iy_start * nx, nx, MPI_REAL_TYPE, top, 0,
                      a_new + (iy_end * nx), nx, MPI_REAL_TYPE, bottom, 0, MPI_COMM_WORLD,
                      MPI_STATUS_IGNORE));
MPI_CALL(MPI_Sendrecv(a_new + (iy_end - 1) * nx, nx, MPI_REAL_TYPE, bottom, 0, a_new, nx,
                      MPI_REAL_TYPE, top, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE));
POP_RANGE
```

# Performance

```
Jacobi relaxation: 1000 iterations on 16384 x 16384 mesh with norm check every 1 iterations
  0, 31.999022
 100, 0.897983
 200, 0.535685
 300, 0.395651
 400, 0.319039
 500, 0.269961
 600, 0.235510
 700, 0.209829
 800, 0.189854
 900, 0.173818
Num GPUs: 4.
16384x16384: 1 GPU: 6.3126 s, 4 GPUs: 1.9474 s, speedup: 3.24, efficiency: 81.04
```


# Profiling

# Instructions

- Install NVIDIA nvidia systems on your laptop
- Run «source /home/guest/init-hpc.sh»
- Run «make profile» to generate nsys profile files
- Copy back the profile files from the cluster to your laptop, and open them with nvidia systems

# Profile


Kernel  
Execution



MPI  
Communication

# Intuition

- Separate computation of top/bottom borders from the inner part
- In this way, we can send/recv immediately the borders while we compute the inner part
- So we would like to have (in parallel):  
`compute_top_kernel<<<...>>>`  
`compute_bottom_kernel<<<...>>>`  
`compute_inner_kernel<<<...>>>`
- Then:
  - wait for the top kernel to finish
  - sendrecv the top border
  - wait for the bottom kernel to finish
  - sendrecv the bottom border
- Issues:
  - By default, each kernel won't start before the termination of the previous one
  - `cudaDeviceSynchronize` waits for all the running kernels to finish (not only for the top/bottom one)



# CUDA Streams


(Let's forget about Multi-GPU for a moment and let's get back to a single GPU scenario)

# PARALLELISM IN SYSTEM ARCHITECTURE


Typical GPU based systems can do multiple things at the same time:

1. Execute a grid on the GPU
2. Copy data from the host to the device
3. Copy data from the device to the host

Max dimension size or a grid size (x,y,z):	(2147483647, 65535, 65535)
Maximum memory pitch:	2147483647 bytes
Texture alignment:	512 bytes
Concurrent copy and kernel execution:	Yes with 2 copy engine(s)
Run time limit on kernels:	No




# VECTOR SCALING SYSTEM UTILIZATION



$$T = T_{H2D} + T_{COMP} + T_{D2H}$$

Q: how to leverage system parallelism?

# PIPELINING



Q: before the execution time was given by  $T_{H2D} + T_{COMP} + T_{D2H}$ . Assuming no extra overhead is introduced with pipelining, what is the now the execution time?

A:  $T_{H2D} + (1/4)*T_{COMP} + (1/4)*T_{D2H}$

# CUDA STREAMS AND ASYNCHRONOUS MEMCPY

CUDA **Streams** support managed concurrent execution of CUDA API functions (e.g., kernel execution and Memcpy-s). They are FIFO command queues:

- ▷ Operations (tasks) in the same stream are serialized
- ▷ Operations (tasks) in different streams are executed in parallel

If no stream is specified, tasks go in the **default stream**. Streams allow you to implement a sort of **task parallelism**.




**Out-of-order** execution for tasks from different streams

To overlap memory copies with host execution, we need to use **asynchronous memory copies**. Pinned memory is necessary for performance reasons.

(Check PMC slides on why/how pinned memory is used)


# Streams

- Requests made from the host code are put into First-In-First-Out queues
  - Queues are read and processed asynchronously by the driver and device
  - Driver ensures that commands in a queue are processed in sequence. E.g., Memory copies end before kernel launch, etc.



# Streams

- To allow concurrent copying and kernel execution, use multiple queues, called “streams”



# CUDA STREAMS AND ASYNCHRONOUS MEMCPY API

Streams have type `cudaStream_t`. To create/destroy/syncronize a stream:

```
cudaError_t cudaStreamCreate(cudaStream_t* pStream);  
cudaError_t cudaStreamDestroy(cudaStream_t Stream);  
cudaError_t cudaStreamSynchronize(cudaStream_t Stream);
```

To create an asynchronous copy in a stream:

```
cudaError_t cudaMemcpyAsync( void* dst, const void* src, size_t count,  
                           cudaMemcpyKind kind, cudaStream_t stream = 0)
```


To memset on a stream:

`cudaMemsetAsync`: same as `cudaMemset`, with the addition of a stream parameter

To launch a grid on a specific stream:

```
my_kernel<<<grid_dim, block_dim, 0, stream>>>(...)
```

# PIPELINING



Q: before the execution time was given by  $T_{H2D} + T_{COMP} + T_{D2H}$ . Assuming no extra overhead is introduced with pipelining, what is the now the execution time?

A:  $T_{H2D} + (1/4)*T_{COMP} + (1/4)*T_{D2H}$

# THE CODE

```
// Kernel code Unchanged ...  
  
int main(int argc, char **argv){  
    float *x_h, *y_, *x_d, *y_d;  
    // Allocate pinned and device memory ...  
  
}
```

# THE CODE

```
// Kernel code Unchanged ...

int main(int argc, char **argv){
    float *x_h, *y_, *x_d, *y_d;
    // Allocate pinned and device memory ...

    const unsigned int num_streams = 32;
    cudaStream_t streams[num_streams];
    for (int i = 0; i < num_streams; i++)
        cudaStreamCreate(&streams[i]);

}

46
```

# THE CODE

```
// Kernel code Unchanged ...

int main(int argc, char **argv){
    float *x_h, *y_, *x_d, *y_d;
    // Allocate pinned and device memory ...

    const unsigned int num_streams = 32;
    cudaStream_t streams[num_streams];
    for (int i = 0; i < num_streams; i++)
        cudaStreamCreate(&streams[i]);

    unsigned int num_segments = num_streams;
    unsigned int segment_size = (N + num_segments - 1) / num_segments;

}
```

# THE CODE

```
// Kernel code Unchanged ...

int main(int argc, char **argv){
    float *x_h, *y_, *x_d, *y_d;
    // Allocate pinned and device memory ...

    const unsigned int num_streams = 32;
    cudaStream_t streams[num_streams];
    for (int i = 0; i < num_streams; i++)
        cudaStreamCreate(&streams[i]);

    unsigned int num_segments = num_streams;
    unsigned int segment_size = (N + num_segments - 1) / num_segments;

    for (int i = 0; i < num_segments; i++) {
        // ...
    }
}
```

# THE CODE

```
// Kernel code Unchanged ...

int main(int argc, char **argv){
    float *x_h, *y_, *x_d, *y_d;
    // Allocate pinned and device memory ...

    const unsigned int num_streams = 32;
    cudaStream_t streams[num_streams];
    for (int i = 0; i < num_streams; i++)
        cudaStreamCreate(&streams[i]);

    unsigned int num_segments = num_streams;
    unsigned int segment_size = (N + num_segments - 1) / num_segments;

    for (int i = 0; i < num_segments; i++) {
        unsigned int start = i * segment_size;
        unsigned int end = (start + segment_size) < N ? start + segment_size : N;
        unsigned int this_segment_size = end - start;

    }
}
```

# THE CODE

```
// Kernel code Unchanged ...

int main(int argc, char **argv){
    float *x_h, *y_, *x_d, *y_d;
    // Allocate pinned and device memory ...

    const unsigned int num_streams = 32;
    cudaStream_t streams[num_streams];
    for (int i = 0; i < num_streams; i++)
        cudaStreamCreate(&streams[i]);

    unsigned int num_segments = num_streams;
    unsigned int segment_size = (N + num_segments - 1) / num_segments;

    for (int i = 0; i < num_segments; i++) {
        unsigned int start = i * segment_size;
        unsigned int end = (start + segment_size) < N ? start + segment_size : N;
        unsigned int this_segment_size = end - start;

        // memcpy async
        cudaMemcpyAsync(x_d + start, x_h + start, this_segment_size * sizeof(float), cudaMemcpyHostToDevice, streams[i]);

        scal<<<(segment_size + 128-1)/128, 128, 0, streams[i]>>>(this_segment_size, alpha, x, y);

        // memcpy async
        cudaMemcpyAsync(y_h + start, y_d + start, this_segment_size * sizeof(float), cudaMemcpyDeviceToHost, streams[i]);
    }

}
```

# THE CODE

```
// Kernel code Unchanged ...

int main(int argc, char **argv){
    float *x_h, *y_, *x_d, *y_d;
    // Allocate pinned and device memory ...

    const unsigned int num_streams = 32;
    cudaStream_t streams[num_streams];
    for (int i = 0; i < num_streams; i++)
        cudaStreamCreate(&streams[i]);

    unsigned int num_segments = num_streams;
    unsigned int segment_size = (N + num_segments - 1) / num_segments;

    for (int i = 0; i < num_segments; i++) {
        unsigned int start = i * segment_size;
        unsigned int end = (start + segment_size) < N ? start + segment_size : N;
        unsigned int this_segment_size = end - start;

        // memcpy async
        cudaMemcpyAsync(x_d + start, x_h + start, this_segment_size * sizeof(float), cudaMemcpyHostToDevice, streams[i]);

        scal<<<(segment_size + 128-1)/128, 128, 0, streams[i]>>>(this_segment_size, alpha, x, y);

        // memcpy async
        cudaMemcpyAsync(y_h + start, y_d + start, this_segment_size * sizeof(float), cudaMemcpyDeviceToHost, streams[i]);
    }
    cudaDeviceSynchronize();
    // Free memory, destroy streams...
}
```


# OVERLAPPING COMMUNICATION AND COMPUTATION

How small the segments should be?

There are filling/emptying phases, slowest stage dominates

We could be tempted to make very fine grained segments, but:

- At a certain point the compute will be too small to fully utilize the GPU
- You have overheads (kernel launching time, DMA setup, ...)



So, have “**moderately**” sized blocks

We can use **pinned memory** and **streams** to overlap communication and computation

- ▷ This requires changing the code
- ▷ Be careful with the use of pinned memory
- ▷ Properly size the segments




# DEFAULT STREAM

- Unless otherwise specified all calls are placed into a default stream
  - Often referred to as “Stream 0”
- Stream 0 has special synchronization rules
  - Synchronous with all streams
    - Operations in stream 0 cannot overlap other streams
- Exception: Streams with non-blocking flag set
  - `cudaStreamCreateWithFlags (&stream, cudaStreamNonBlocking)`
  - Use to get concurrency with libraries out of your control (e.g. MPI)

# KERNEL CONCURRENCY


- Assume foo only utilizes 50% of the GPU
- Default stream

```
foo<<<blocks, threads>>>();  
foo<<<blocks, threads>>>();
```



- Default & user streams

```
cudaStream_t stream1;  
cudaStreamCreate(&stream1);  
foo<<<blocks, threads>>>();  
foo<<<blocks, threads, 0, stream1>>>();  
cudaStreamDestroy(stream1);
```






## KERNEL CONCURRENCY

- Assume foo only utilizes 50% of the GPU
- Default & user streams


```
cudaStream_t stream1;  
cudaStreamCreateWithFlags(&stream1, cudaStreamNonBlocking);  
foo<<<blocks, threads>>>();  
foo<<<blocks, threads, 0, stream1>>>();  
cudaStreamDestroy(stream1);
```



# Back to Profiling and Multi GPU

# Profile


Kernel  
Execution



MPI  
Communication


# Expectations

ParaStationMPI 5.4.10-1 – JUWELS Booster – NVIDIA A100 40 GB – Jacobi on 17408x17408




# Intuition

- Separate computation of top/bottom borders from the inner part
- In this way, we can send/recv immediately the borders while we compute the inner part
- So we would like to have (in parallel):  
`compute_top_kernel<<<...>>>`  
`compute_bottom_kernel<<<...>>>`  
`compute_inner_kernel<<<...>>>`
- Then:
  - wait for the top kernel to finish
  - sendrecv the top border
  - wait for the bottom kernel to finish
  - sendrecv the bottom border
- Issues:
  - By default, each kernel won't start before the termination of the previous one
  - `cudaDeviceSynchronize` waits for all the running kernels to finish (not only for the top/bottom one)



# Communication + Computation Overlap



# MPI Communication + Computation Overlap

```
launch_jacobi_kernel(a_new,      a,  l2_norm_d,  iy_start,  (iy_start + 1),  nx,  push_top_stream);
launch_jacobi_kernel(a_new,      a,  l2_norm_d,  (iy_end - 1), iy_end,  nx,  push_bottom_stream);
launch_jacobi_kernel(a_new,      a,  l2_norm_d,  (iy_start + 1), (iy_end - 1), nx,  compute_stream);

const int top = rank > 0 ? rank - 1 : (size - 1);
const int bottom = (rank + 1) % size;

CUDA_RT_CALL(cudaStreamSynchronize(push_top_stream));
MPI_CALL(MPI_Sendrecv(a_new + iy_start * nx, nx, MPI_REAL_TYPE, top, 0,
                      a_new + (iy_end * nx), nx, MPI_REAL_TYPE, bottom, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE));
CUDA_RT_CALL(cudaStreamSynchronize(push_bottom_stream));
MPI_CALL(MPI_Sendrecv(a_new + (iy_end - 1) * nx, nx, MPI_REAL_TYPE, bottom, 0,
                      a_new,                  nx, MPI_REAL_TYPE, top, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE));
```

# Race Conditions with Streams

```
while (l2_norm > tol && iter < iter_max) {  
    CUDA_RT_CALL(cudaMemsetAsync(l2_norm_d, 0, sizeof(real), compute_stream));  
  
    calculate_norm = (iter % nccheck) == 0 || (!csv && (iter % 100) == 0);  
  
    launch_jacobi_kernel(a_new, a, l2_norm_d, iy_start, iy_end, nx, calculate_norm,  
                         compute_stream);
```


- At the beginning of every iteration, `l2_norm_d` is set to 0. This is done with a `cudaMemsetAsync`
- This operation must be run on a specific stream (`compute_stream` in this case)
- If there are other kernels enqueued on other streams (e.g., those for computing top/bottom boundaries), those might start before this `cudaMemset` is completed, and would work on a wrong `l2_norm_d` value
- Thus, we must ensure that those kernels running on other streams, do not start before the completion of this `cudaMemsetAsync`

# CUDA Events

(Let's forget about Multi-GPU for a moment and let's get back to a single GPU scenario)

# CUDA Events

- CUDA “events” allow the host thread to query and synchronize with individual streams
- By checking these event I can know what have been processed so far from the stream queue



# CUDA Events

- Events in CUDA can be used to capture the time instance a GPU operation completes.
- Events allow fine-grained timing of GPU operations.
- Events are represented as instances of the `cudaEvent_t` type and they can be:
  - **Created:**

```
cudaError_t cudaEventCreate (cudaEvent_t *event)
```

or

```
cudaEventCreateWithFlags(&ev, cudaEventDisableTiming)
```

- Disable timing to increase performance

- **Destroyed:**

```
cudaError_t cudaEventDestroy (cudaEvent_t event)
```

- **Captured (i.e., enqueueued on the stream queue):**

```
cudaError_t cudaEventRecord (cudaEvent_t event, cudaStream_t stream=0)
```

- **Waited for:**

```
cudaError_t cudaEventSynchronize (cudaEvent_t event)
```

# CUDA Events

- Once two events have been triggered, (e.g. After `cudaEventSynchronize`, `cudaDeviceSynchronize`, `cudaStreamDestroy`, etc...), the duration between two events can be calculated via:

```
cudaError_t cudaEventElapsedTime(
    float *ms,           // Storage for elapsed time
                        // in msec units (OUT)
    cudaEvent_t start,   // Start event identifier (IN)
    cudaEvent_t end);    // End event identifier (IN)
```

# CUDA Events

E.g., timing host-to-device memory copy:

```
cudaEvent_t startT, endT;
float duration;
cudaEventCreate (&startT);
cudaEventCreate (&endT);

for (dataSize = 0; dataSize <= MAXDATASIZE; dataSize += step)
{
    cudaEventRecord (startT, str);
    for (i = 0; i < iter; i++)
    {
        cudaMemcpyAsync (d_data, h_data, sizeof (int) * dataSize, cudaMemcpyHostToDevice, str);
    }
    cudaEventRecord (endT, str);
    cudaEventSynchronize (endT);
    cudaEventElapsedTime (&duration, startT, endT);
    printf ("%i %f\n", (int) (dataSize * sizeof (int)), duration / iter);
}
```

Testing for different data block sizes

Testing multiple times

# Synchronization using Events

- `cudaEventSynchronize(cudaEvent_t event)`
  - Host wait for the events to occur (i.e., all the preceding operations on the stream completed)
- `cudaEventQuery(cudaEvent_t event)`
  - Host checks if the event occurred (without blocking)
- `cudaStreamWaitEvent(cudaStream_t stream, cudaEvent_t event, unsigned int flags = 0)`
  - The subsequent work submitted to the stream "stream" must wait until the event "event" occurs
  - Does not block the host!
- ATTENTION: do not call `cudaEventSynchronize` (it is blocking!) before enqueueing that event with `cudaEventRecord`

# Synchronization using Events - Recap

Synchronization between host and device through:

- `cudaDeviceSynchronize`
- `cudaStreamSynchronize`
- `cudaEventSyncrhonize`

Synchronization between streams:

- `cudaStreamWaitEvent`

# Recap

- Starting point: having to wait for all the tile to be computed before sending/receiving is inefficient
- Let's separate the computation in three kernels: 1 for bottom boundary, 1 for top boundary, one for inner rows, so that we can overlap computation and communication
- Let's put these three kernels on three different streams
- However, we need to pay attention to race conditions, we might need to synchronize events between different streams (e.g., setting `l2_norm` to 0 before any kernel can start)

# Back to Comp. Comm. Overlap: Hands-On

# Instructions


- Join the assignment on Github Classroom: <https://classroom.github.com/a/K8voN6eF>
  - At the moment, there is no auto-grading yet, check correctness by yourself
- Access the cluster and clone the repository
- Fill the «TODOs» in the jacobi.cu code
- Run “source /home/guest/init-hpc.sh” to load the MPI, CUDA compilers etc...
- Run «make run» to compile and run the code (by default on 4 GPUs, one GPU per node)

# Performance

```
Jacobi relaxation: 1000 iterations on 16384 x 16384 mesh with norm check every 1 iterations
  0, 31.999020
 100, 0.897983
 200, 0.535685
 300, 0.395651
 400, 0.319039
 500, 0.269961
 600, 0.235510
 700, 0.209829
 800, 0.189854
 900, 0.173818
Num GPUs: 4.
16384x16384: 1 GPU:  5.9741 s, 4 GPUs:  1.5338 s, speedup: 3.90, efficiency: 97.38
```


# Profile (Before, no overlap)

Kernel  
Execution



MPI  
Communication

# Profile (Now, with overlap)



Less Idle Time than before


# Profile (Now, with overlap)




3 different streams

# Communication + Computation Overlap

ParaStationMPI 5.4.10-1 – JUWELS Booster – NVIDIA A100 40 GB – Jacobi on 17408x17408



# Profile (Now, with overlap)



# MPI Communication + Computation Overlap

```
launch_jacobi_kernel(a_new, a, 12_norm_d, (iy_start + 1), (iy_end - 1), nx, compute_stream);
launch_jacobi_kernel(a_new, a, 12_norm_d, iy_start, (iy_start + 1), nx, push_top_stream);
launch_jacobi_kernel(a_new, a, 12_norm_d, (iy_end - 1), iy_end, nx, push_bottom_stream);
launch_jacobi_kernel(a_new, a, 12_norm_d, (iy_start + 1), (iy_end - 1), nx, compute_stream);

const int top = rank > 0 ? rank - 1 : (size - 1);
const int bottom = (rank + 1) % size;

CUDA_RT_CALL(cudaStreamSynchronize(push_top_stream));
MPI_CALL(MPI_Sendrecv(a_new + iy_start * nx, nx, MPI_REAL_TYPE, top, 0,
                      a_new + (iy_end * nx), nx, MPI_REAL_TYPE, bottom, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE));
CUDA_RT_CALL(cudaStreamSynchronize(push_bottom_stream));
MPI_CALL(MPI_Sendrecv(a_new + (iy_end - 1) * nx, nx, MPI_REAL_TYPE, bottom, 0,
                      a_new, nx, MPI_REAL_TYPE, top, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE));
```




# Performance

```
-----  
Jacobi relaxation: 1000 iterations on 16384 x 16384 mesh with norm check every 1 iterations  
0, 31.999020  
100, 0.897983  
200, 0.535685  
300, 0.395651  
400, 0.319039  
500, 0.269961  
600, 0.235510  
700, 0.209829  
800, 0.189854  
900, 0.173818  
Num GPUs: 4.  
16384x16384: 1 GPU: 5.9750 s, 4 GPUs: 1.7216 s, speedup: 3.47, efficiency: 86.7%
```

Why lower speedup/efficiency?

# Profiling



Even if kernel on different streams can go in parallel, the first kernel to be launched is eating up all the resources and the two smaller kernels can't start before it is over, falling back to the non-overlapping case

# High Priority Streams

Improve scalability with high priority streams (available on CC 3.5+)

```
cudaStreamCreateWithPriority ( cudaStream_t* pStream, unsigned int flags, int priority )
```

- If something else is running with a lower priority on a different stream, hold that and run this first, and then resume it when the stuff on this stream completes

# MPI Communication + Computation Overlap

with high priority streams

```
int leastPriority = 0;
int greatestPriority = leastPriority;

CUDA_RT_CALL(cudaDeviceGetStreamPriorityRange(&leastPriority, &greatestPriority));


cudaStream_t compute_stream;
cudaStream_t push_top_stream;
cudaStream_t push_bottom_stream;

CUDA_RT_CALL(cudaStreamCreateWithPriority(&compute_stream, cudaStreamDefault, leastPriority));
CUDA_RT_CALL(cudaStreamCreateWithPriority(&push_top_stream, cudaStreamDefault, greatestPriority));
CUDA_RT_CALL(cudaStreamCreateWithPriority(&push_bottom_stream, cudaStreamDefault, greatestPriority));
```

# Performance

```
Jacobi relaxation: 1000 iterations on 16384 x 16384 mesh with norm check every 1 iterations
  0, 31.999022
100, 0.897983
200, 0.535685
300, 0.395651
400, 0.319039
500, 0.269961
600, 0.235510
700, 0.209829
800, 0.189854
900, 0.173818
Num GPUs: 4.
16384x16384: 1 GPU: 6.0214 s, 4 GPUs: 1.5456 s, speedup: 3.90, efficiency: 97.40
```


# Profiling



Now the three kernels, in three different streams, run «in parallel» (the smaller kernels might preempt the larger kernel) with the sendrecv

# Communication + Computation Overlap

ParaStationMPI 5.4.10-1 – JUWELS Booster – NVIDIA A100 40 GB – Jacobi on 17408x17408



Source: <https://github.com/NVIDIA/multi-gpu-programming-models>  
JUWELS Booster: <https://apps.fz-juelich.de/jsc/hps/juwels/booster-overview.html>